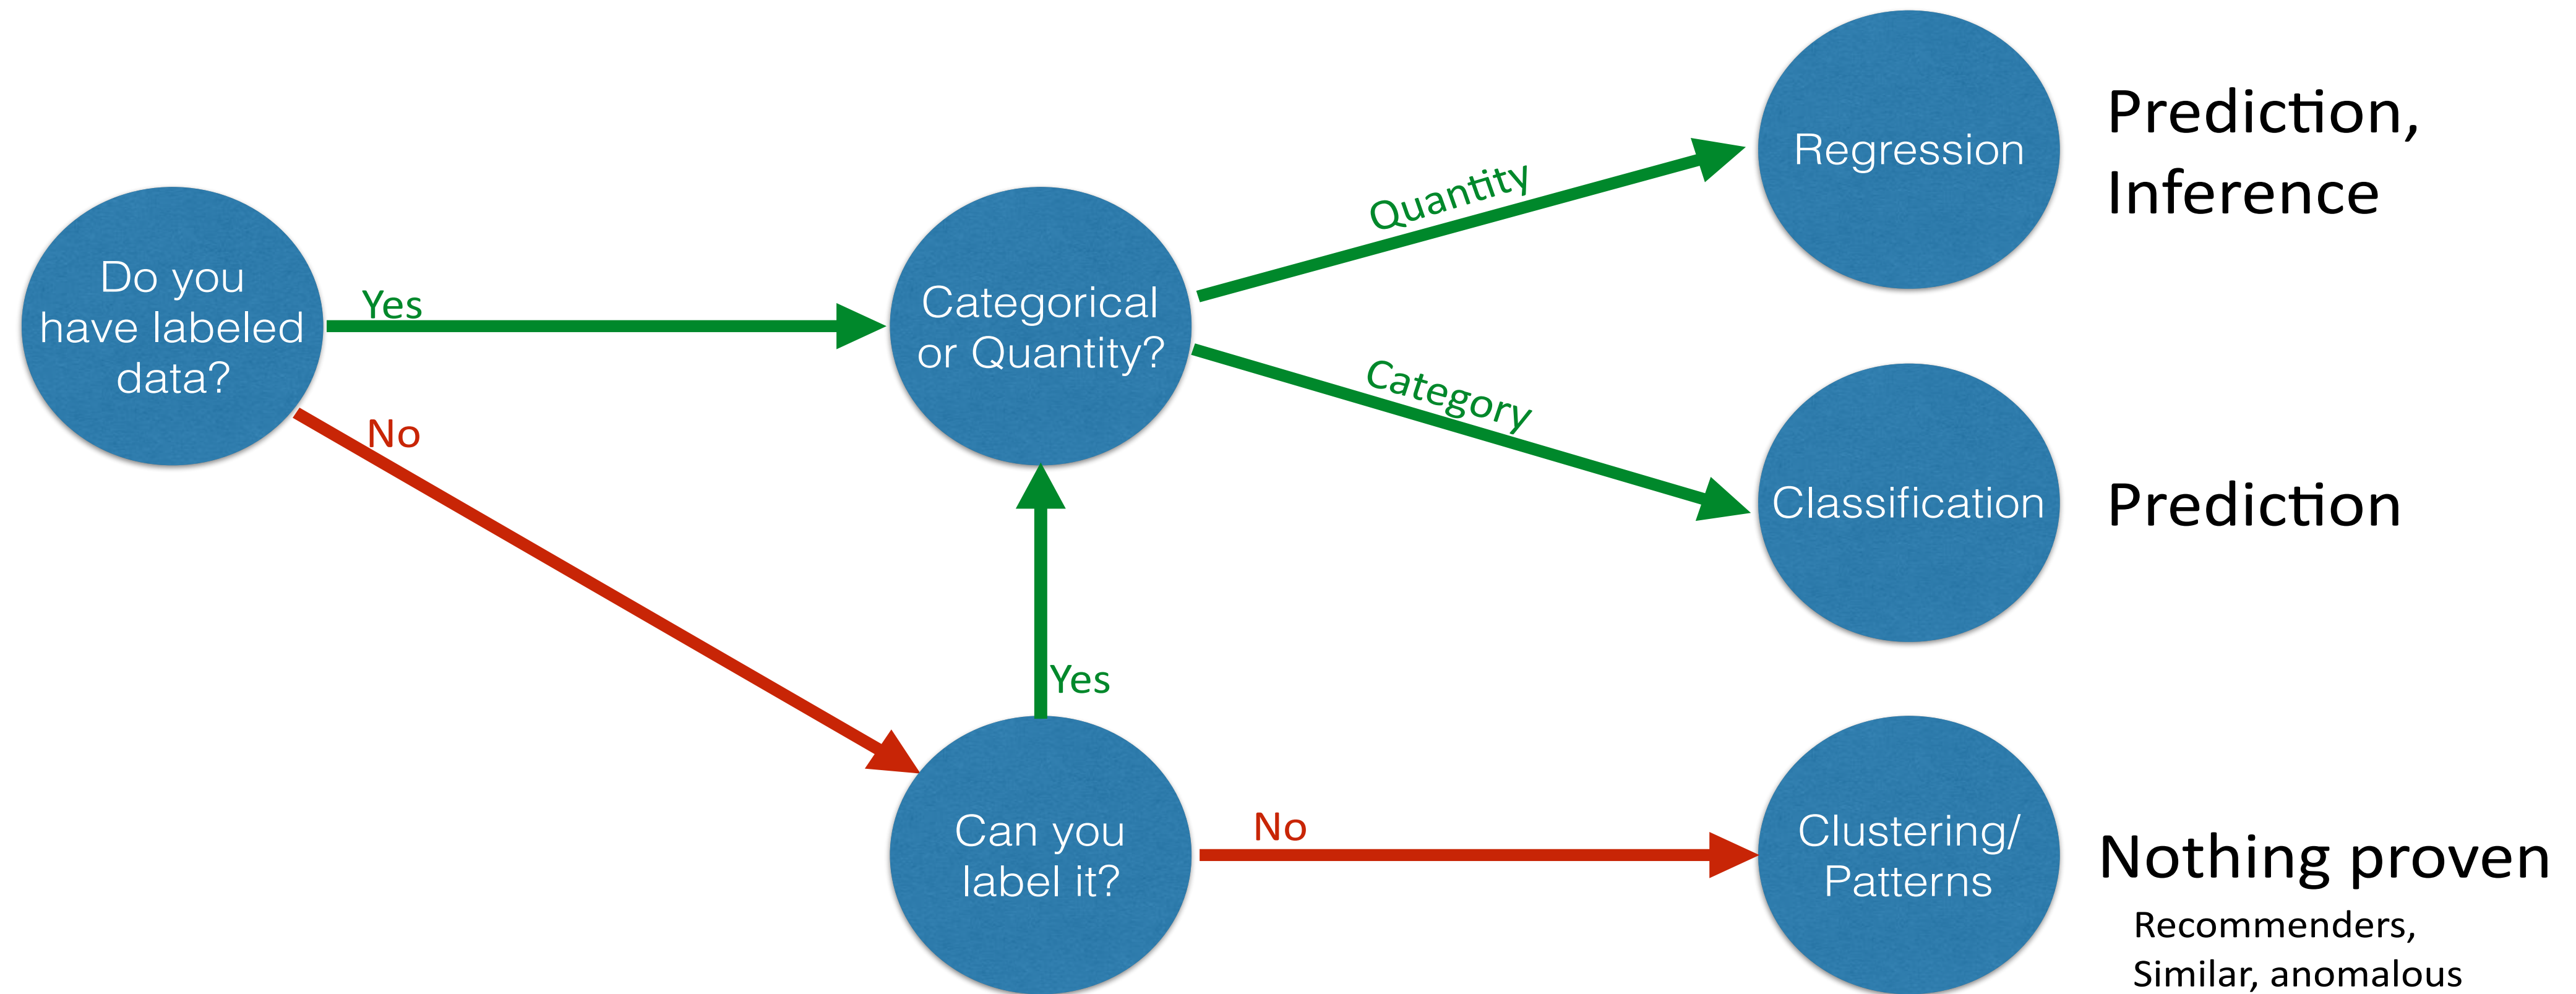


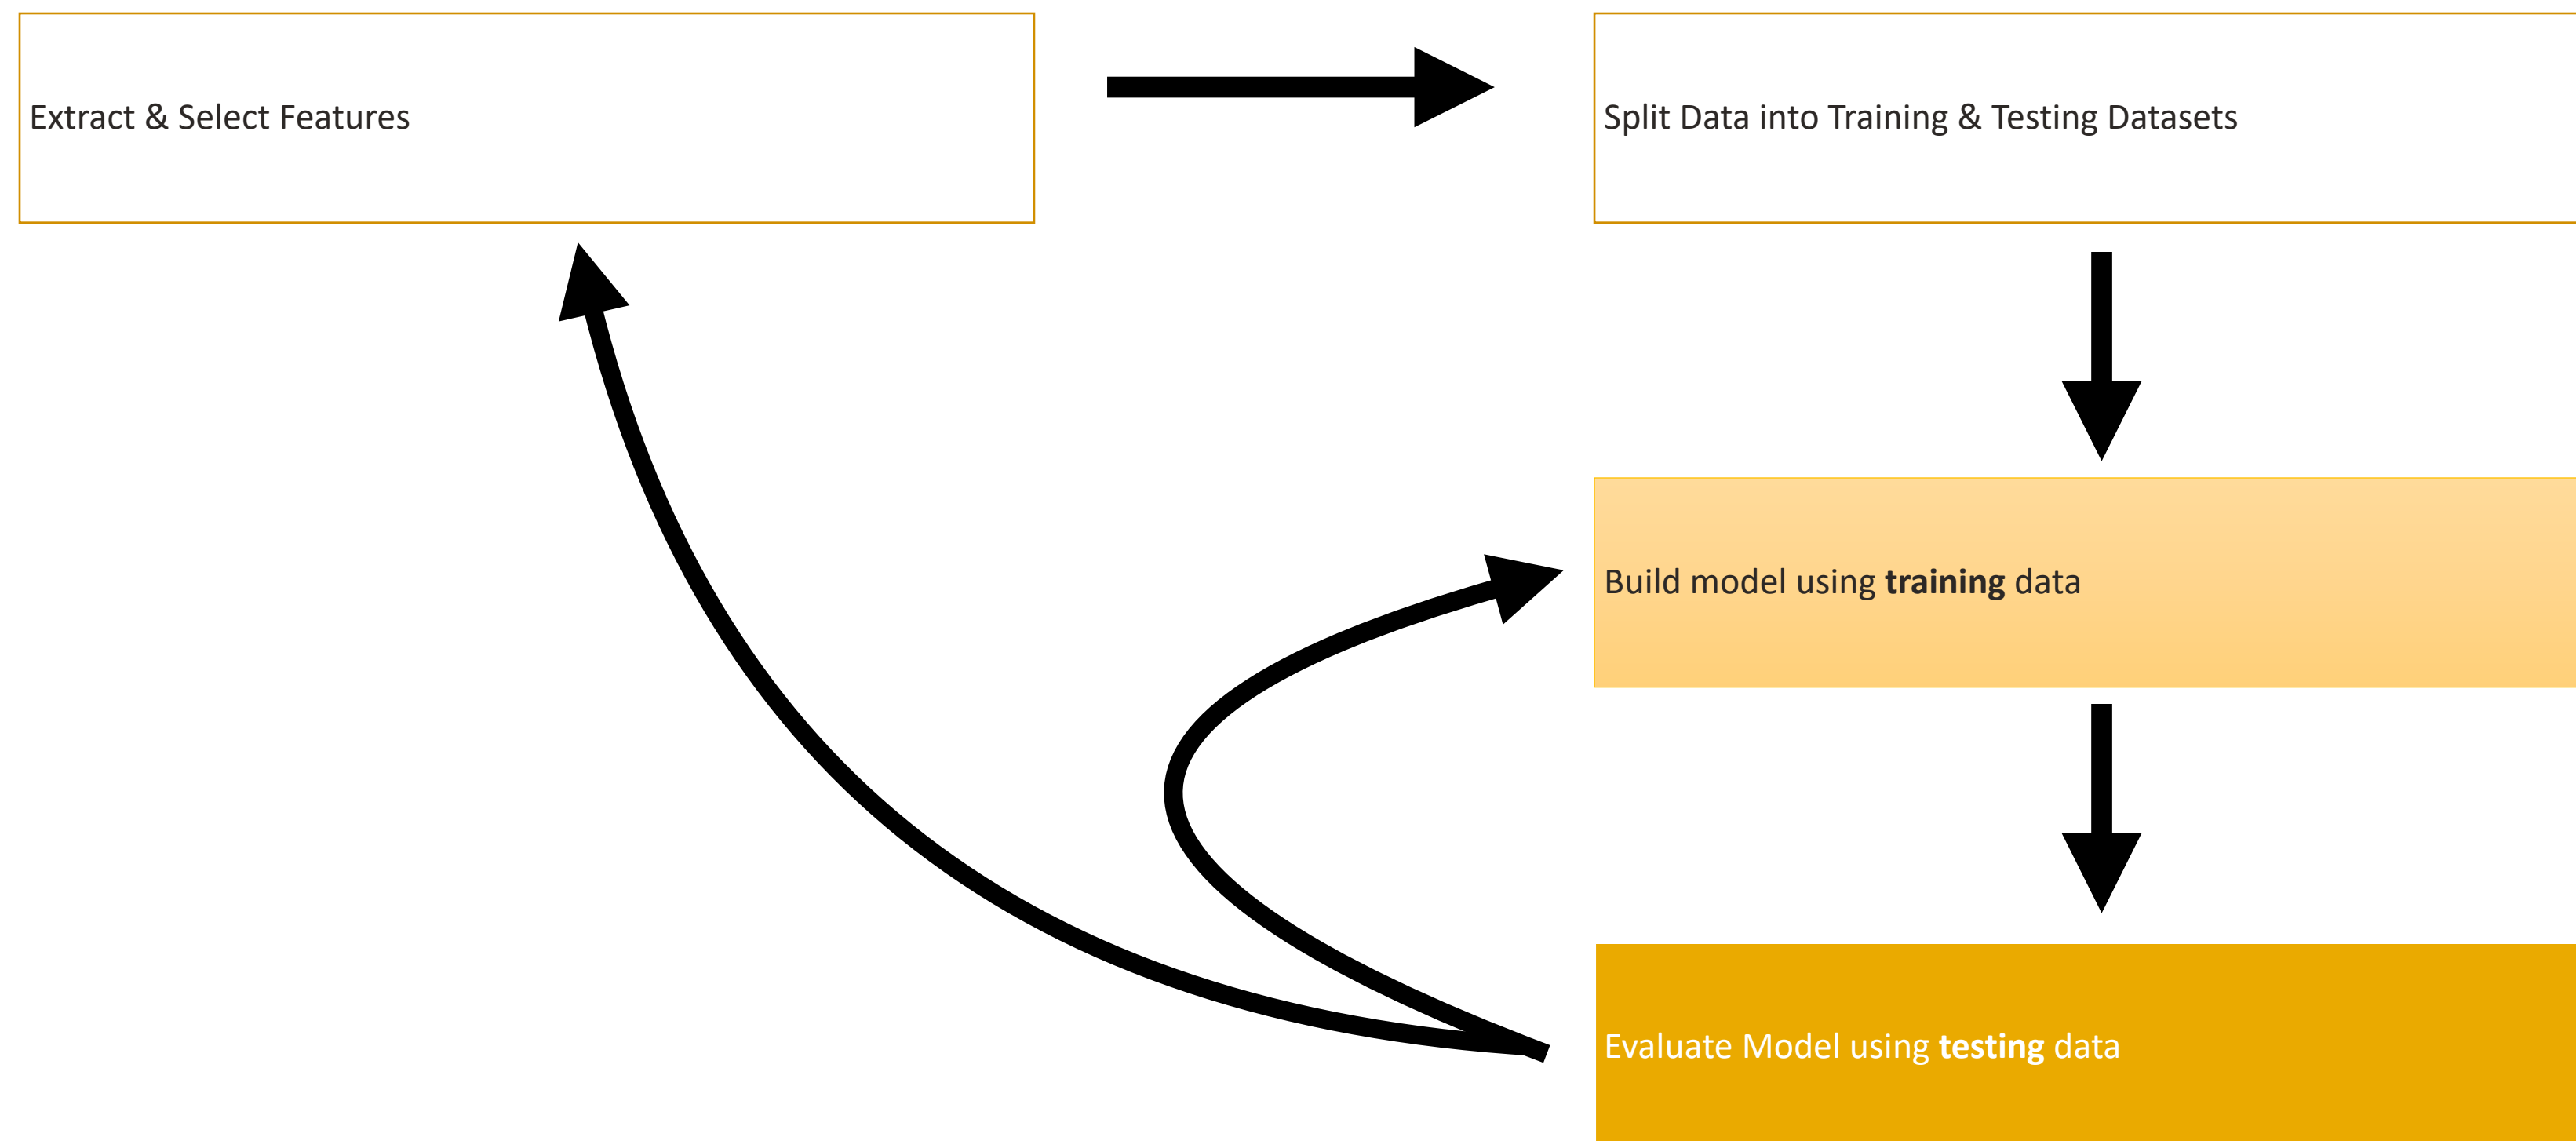
# Module 6: Unsupervised Learning: Clustering

# Agenda for Today

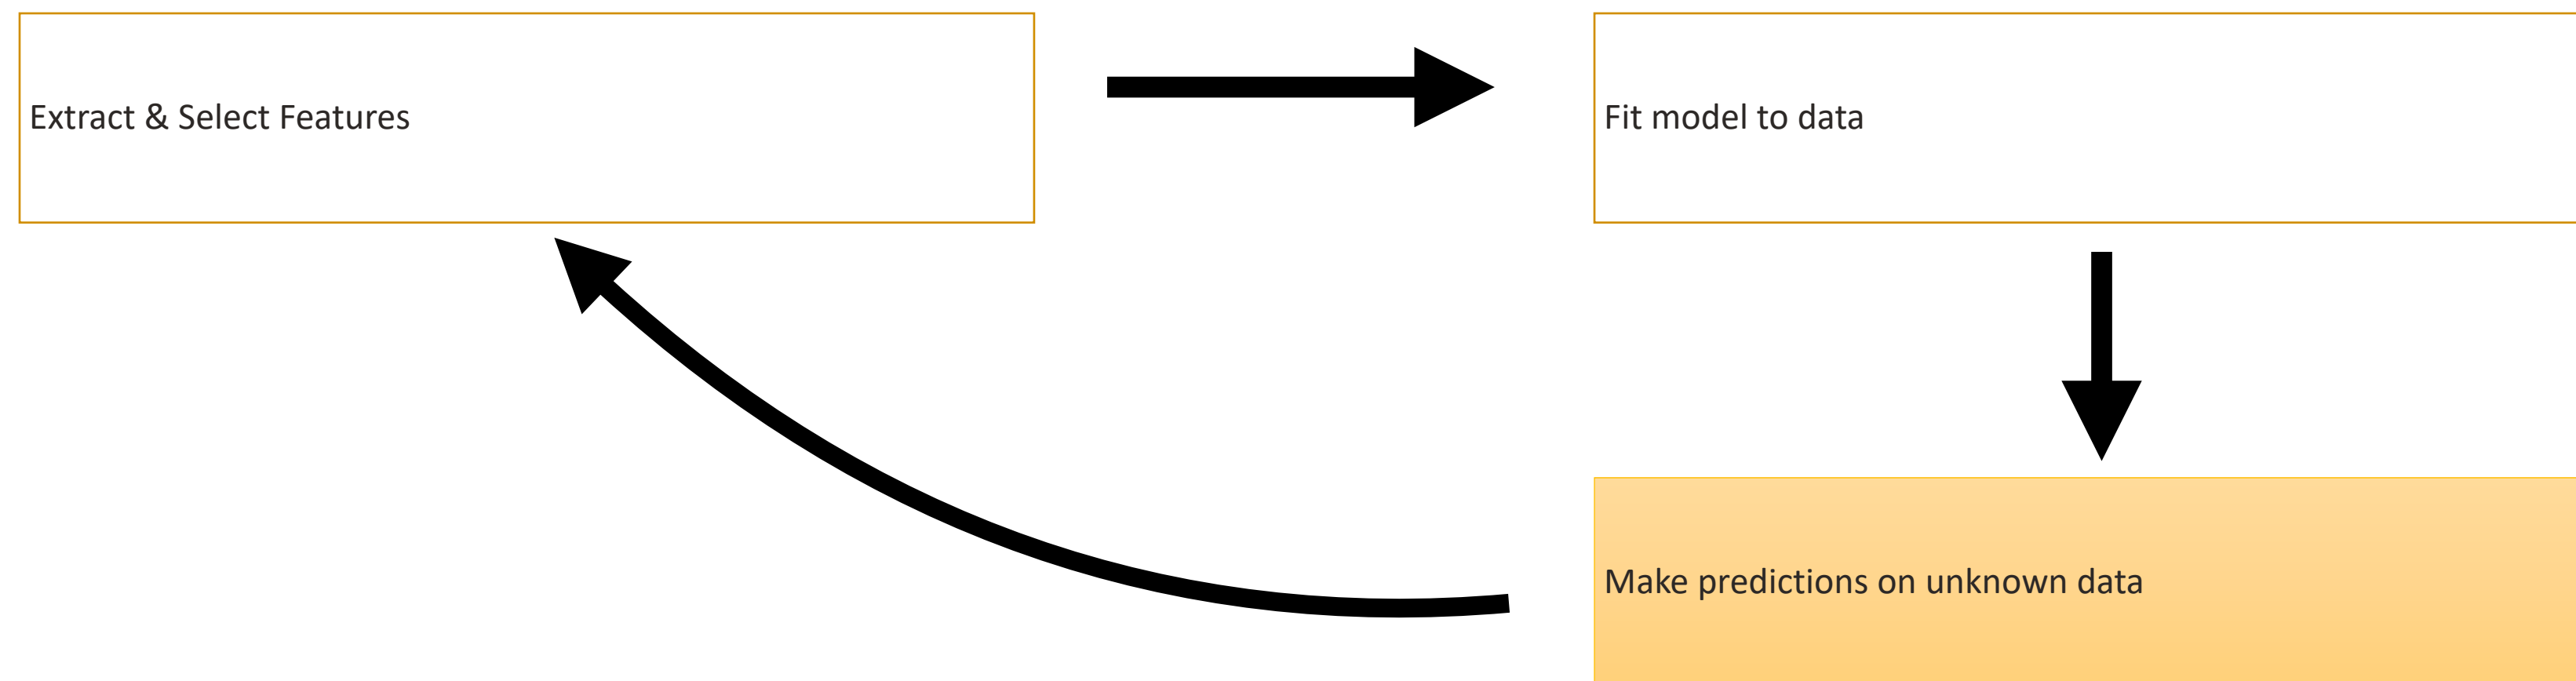
- Measuring Distances
- Math free overview of clustering techniques



# Supervised ML Process



# Unsupervised ML Process



# Unsupervised Clustering Algorithm

1. Select Features
2. Calculate a distance measure
3. Apply a clustering algorithm
4. Validate?

# Which Departments are Similar?

	Malware events
Dept1	6
Dept2	1
Dept3	8

# Which Departments are Similar?

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1





# Which Departments are Similar?

	Malware events	Phishing	Open Tickets
Dept1	6	6	3
Dept2	1	2	1
Dept3	8	1	9

# Computing Distance

	Malware events
Dept1	6
Dept2	1
Dept3	8

Compare:

Dept1 to Dept2:  $|6 - 1| = 5$

Dept2 to Dept3:  $|1 - 8| = 7$

Dept1 to Dept3:  $|6 - 8| = 2$

# Two-Dimensional Distance

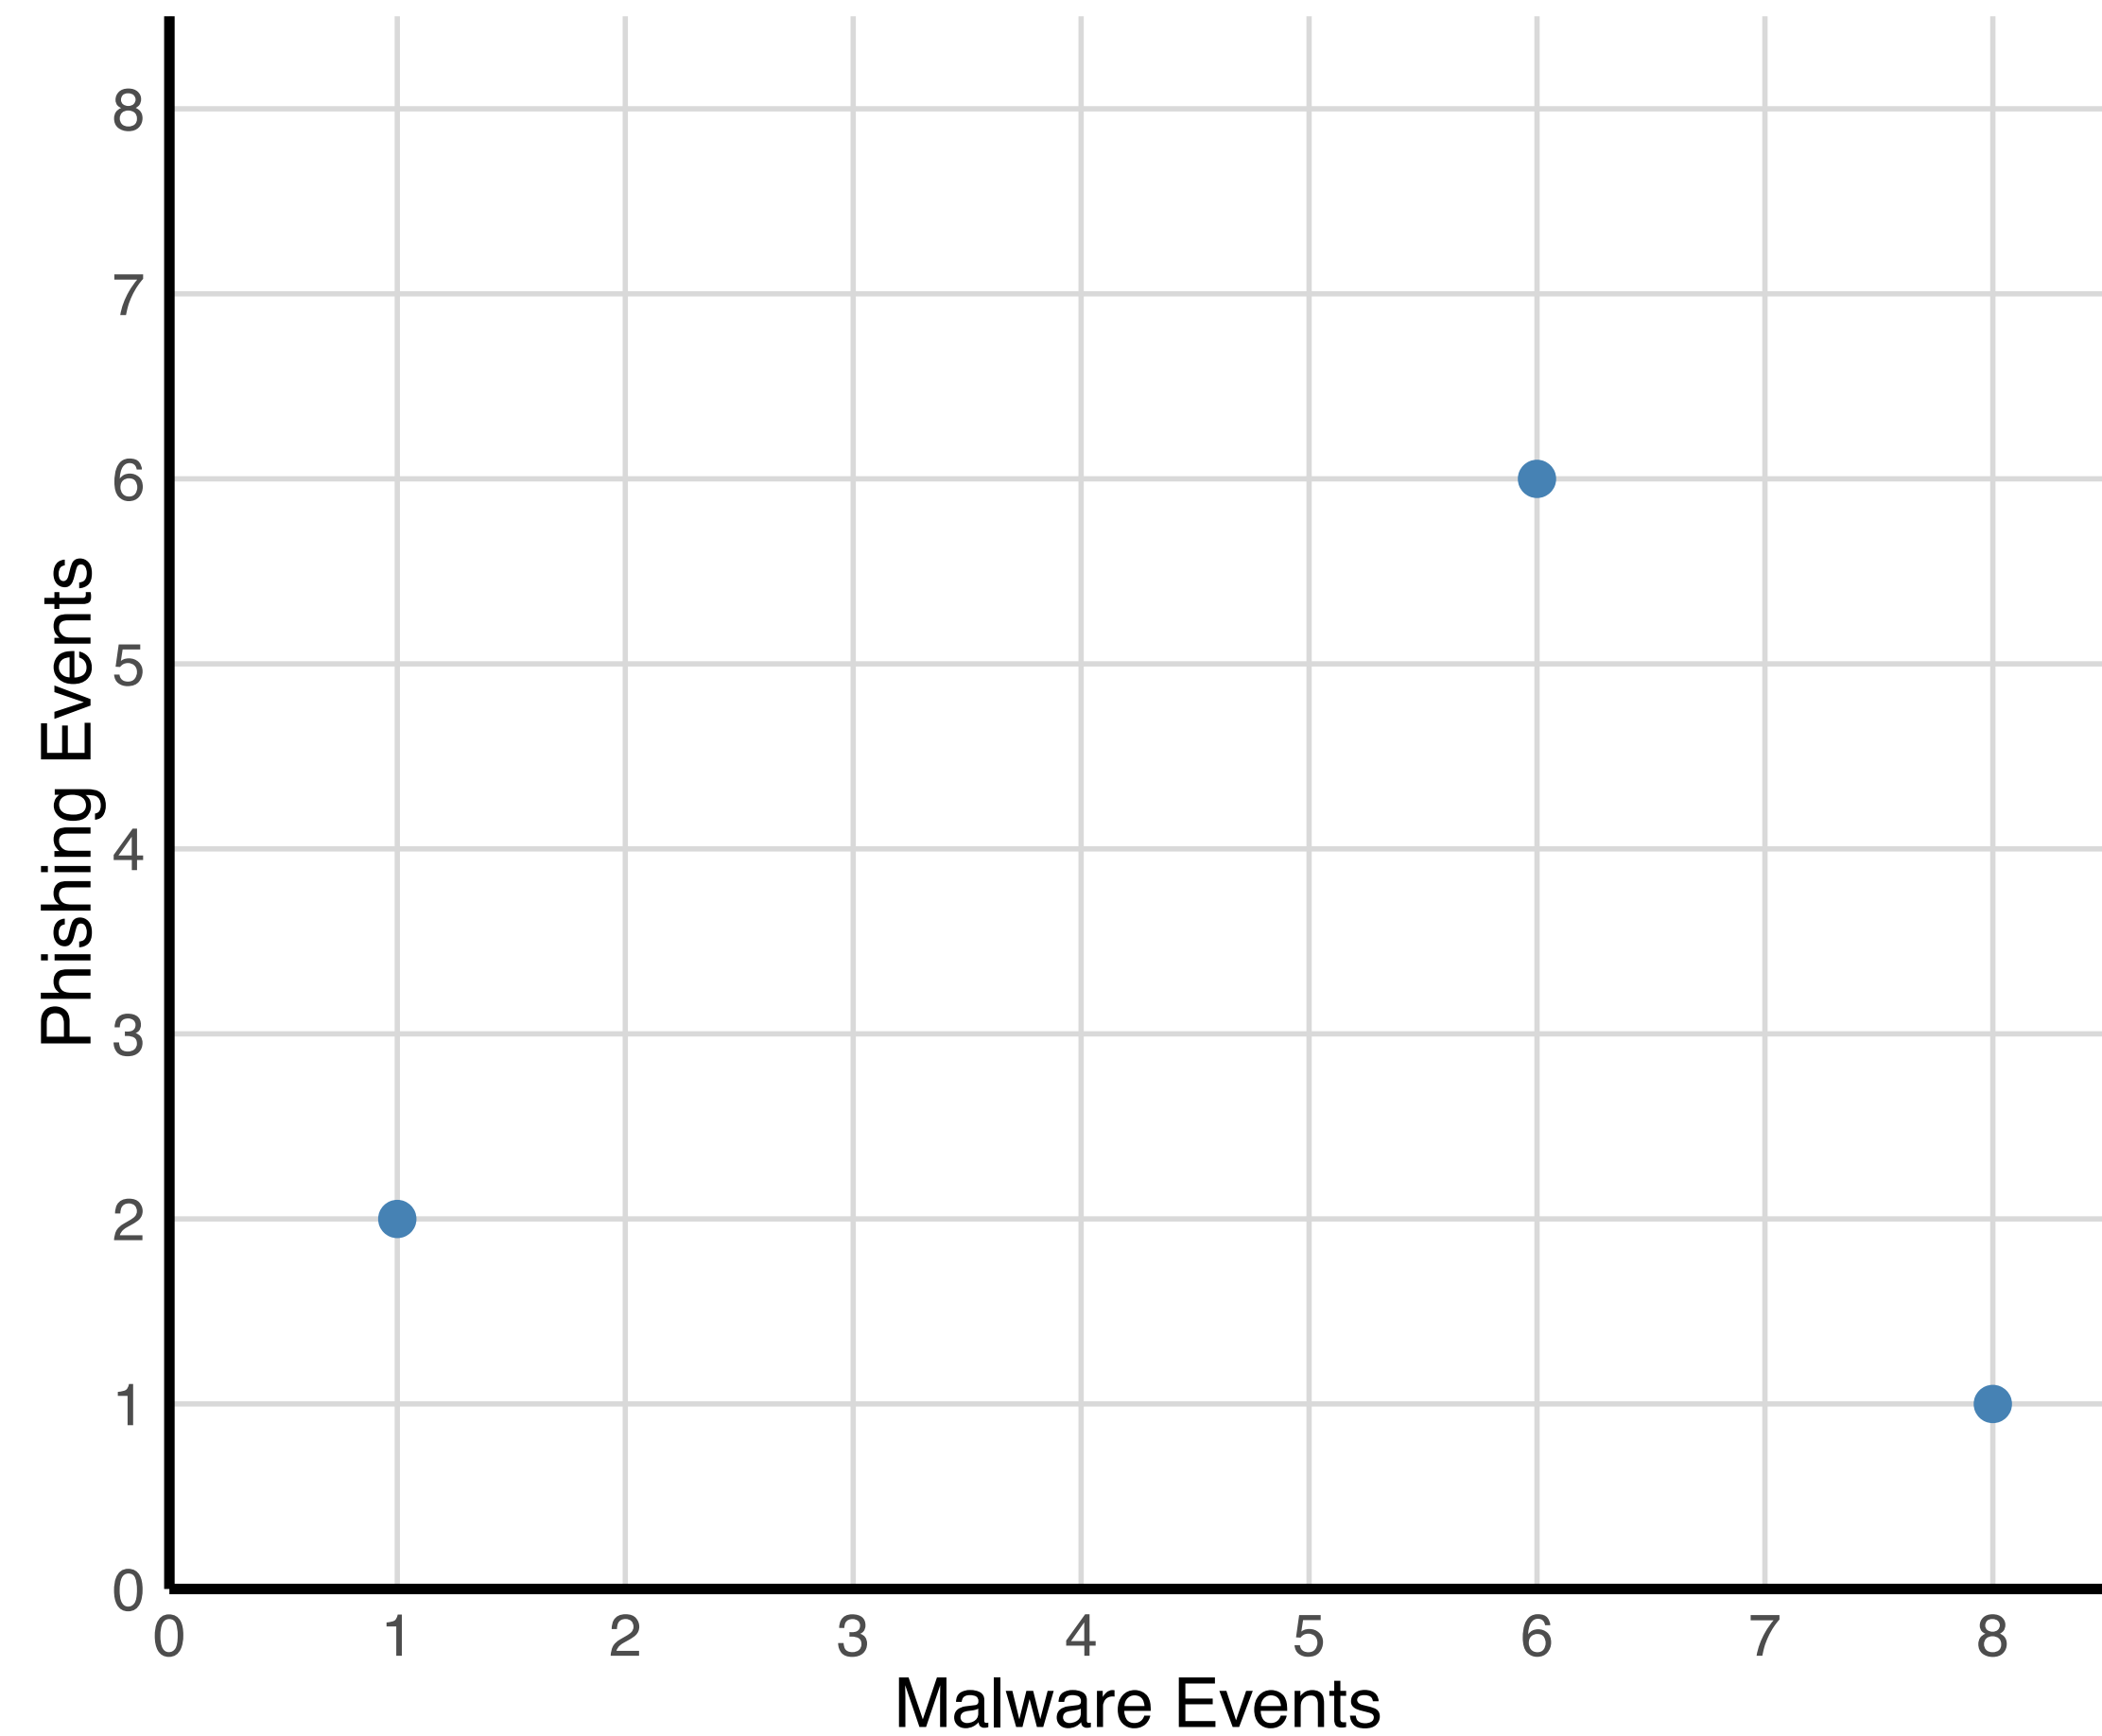
	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1

Multiple Distance methods

- Euclidean
  - Manhattan
  - Maximum
  - Canberra
  - Binary
  - Minkowski
- ... (to name a few)

# Two-Dimensional Distance


	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



# Two-Dimensional Distance

Euclidean very common and easy to understand

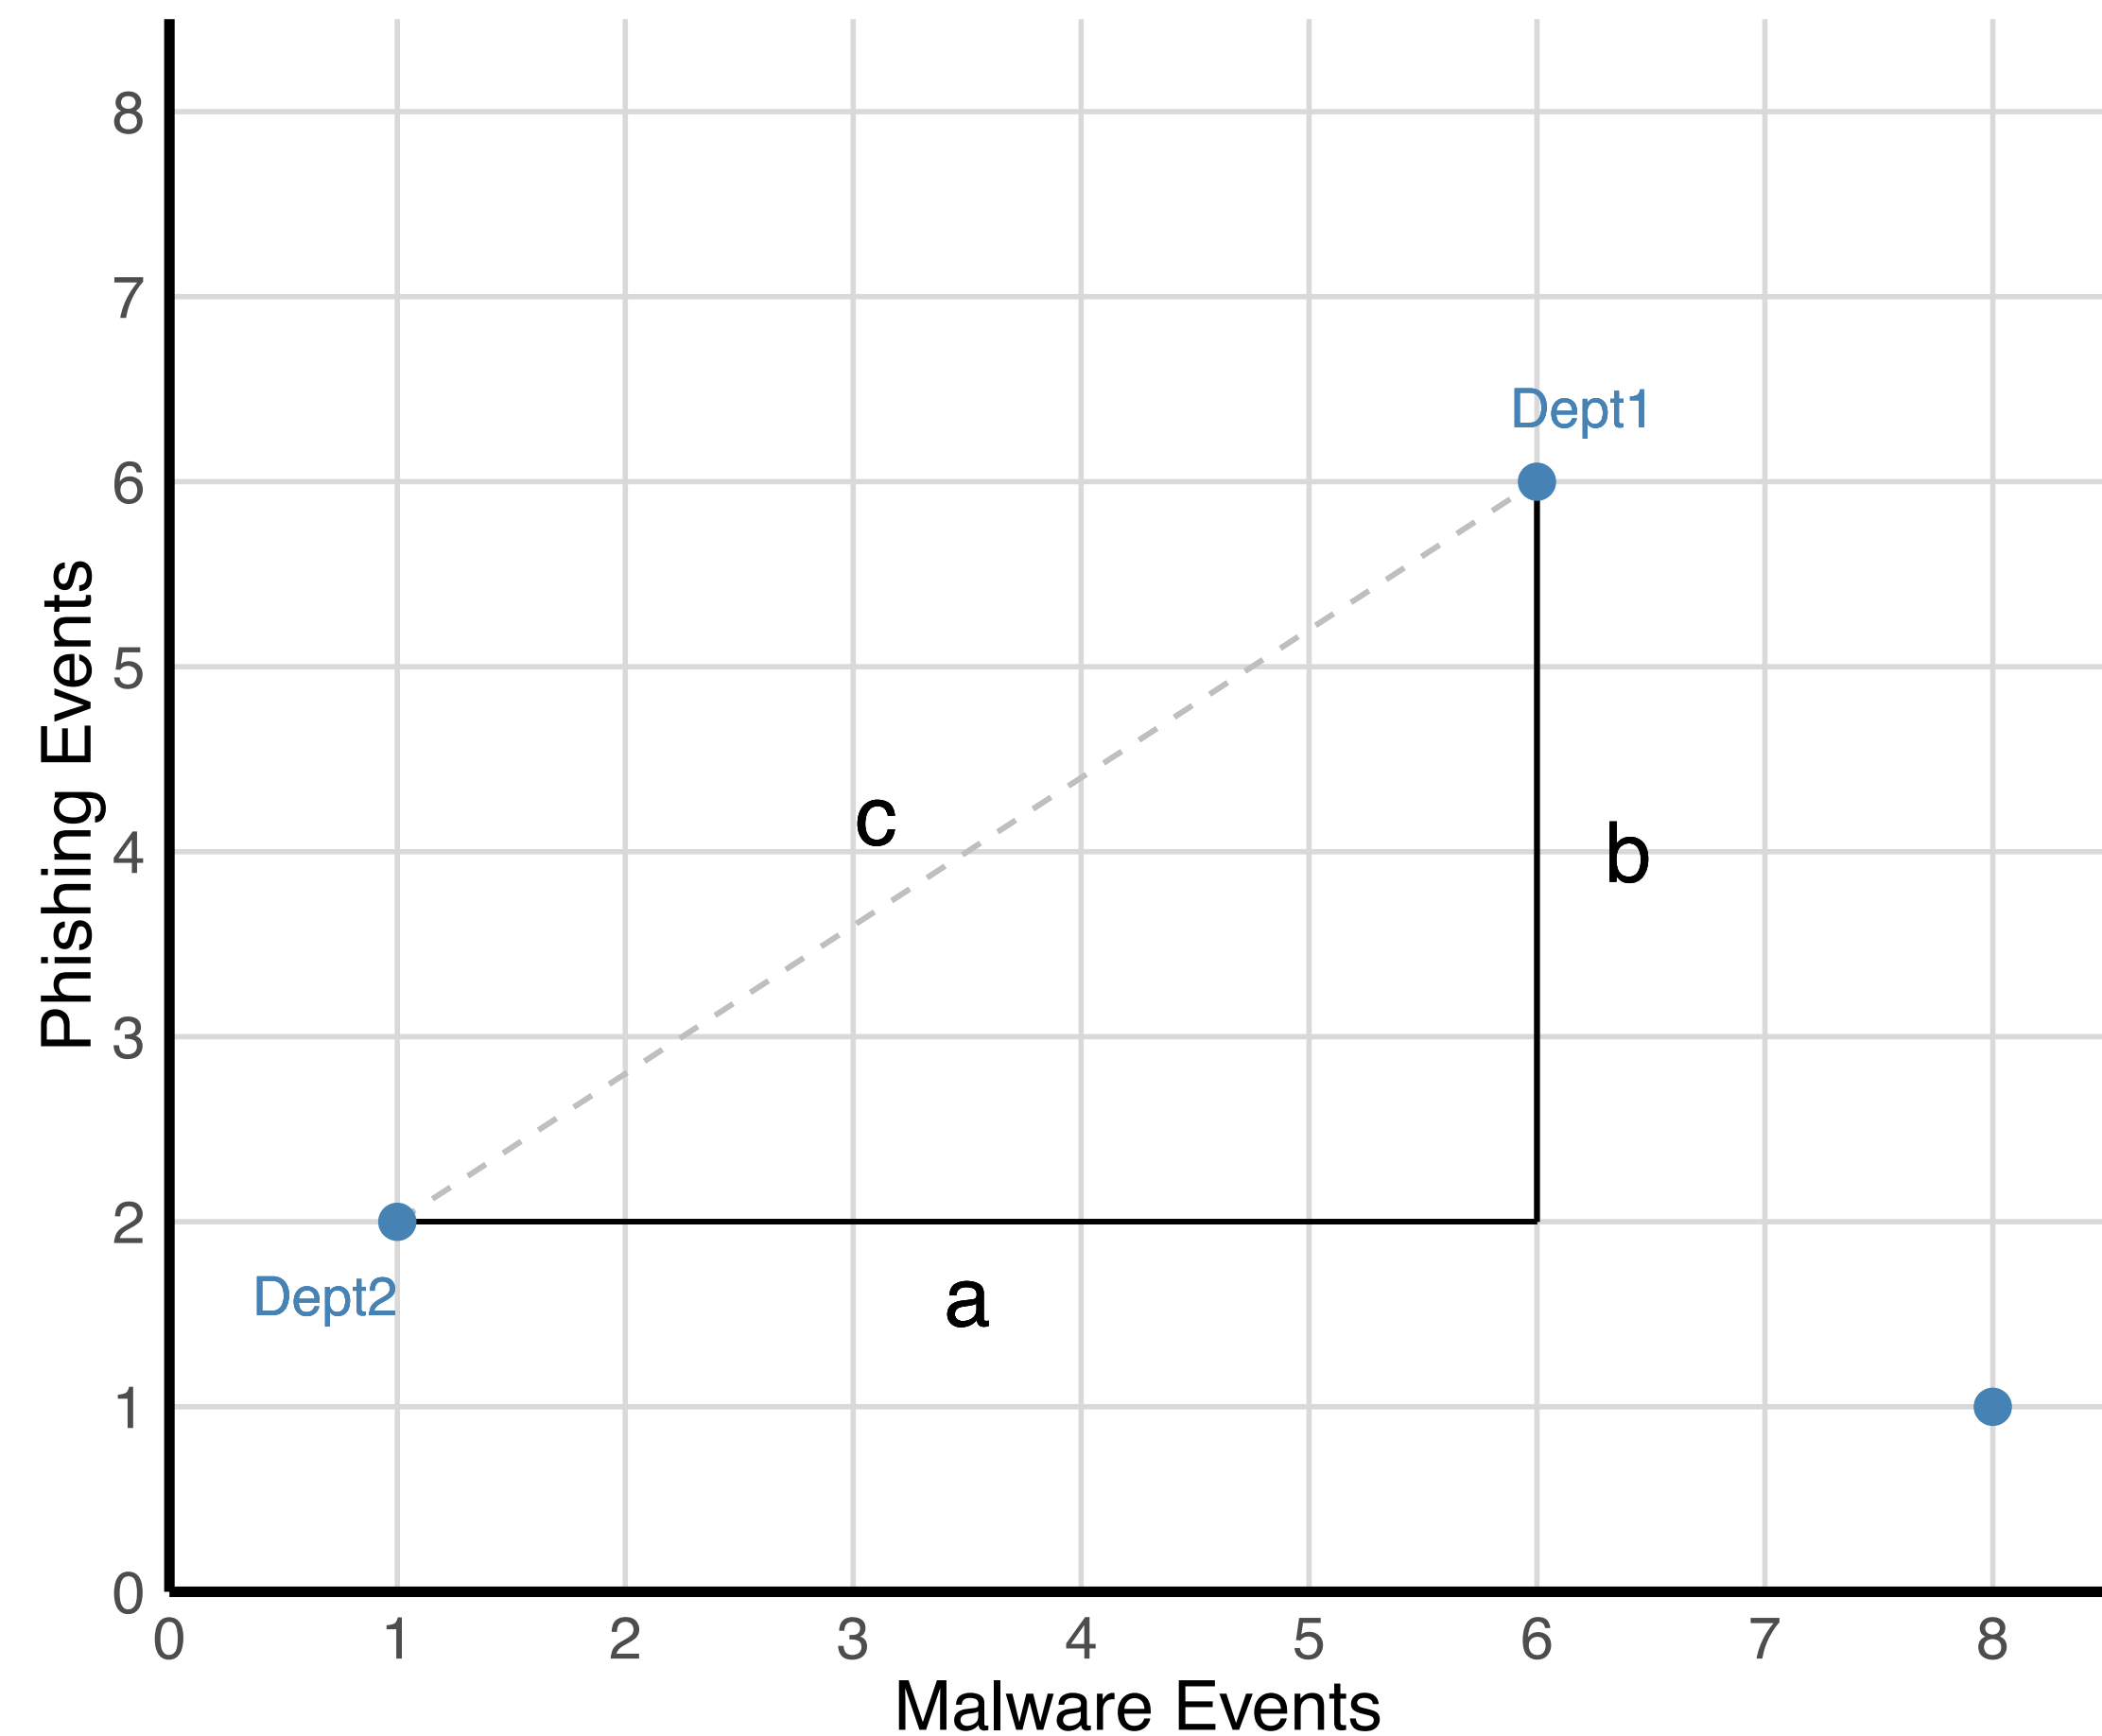
	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



# Two-Dimensional Distance

Euclidean very common and easy to understand:  $a^2 + b^2 = c^2$


	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



# Two-Dimensional Distance

Manhattan also easy to comprehend:  $a + b$

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



# Computing Distance

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1

Compare:

Dept1 to Dept2:  $\sqrt{(6-1)^2 + (6-2)^2} = \mathbf{6.4}$

Dept2 to Dept3: ... = **7.1**

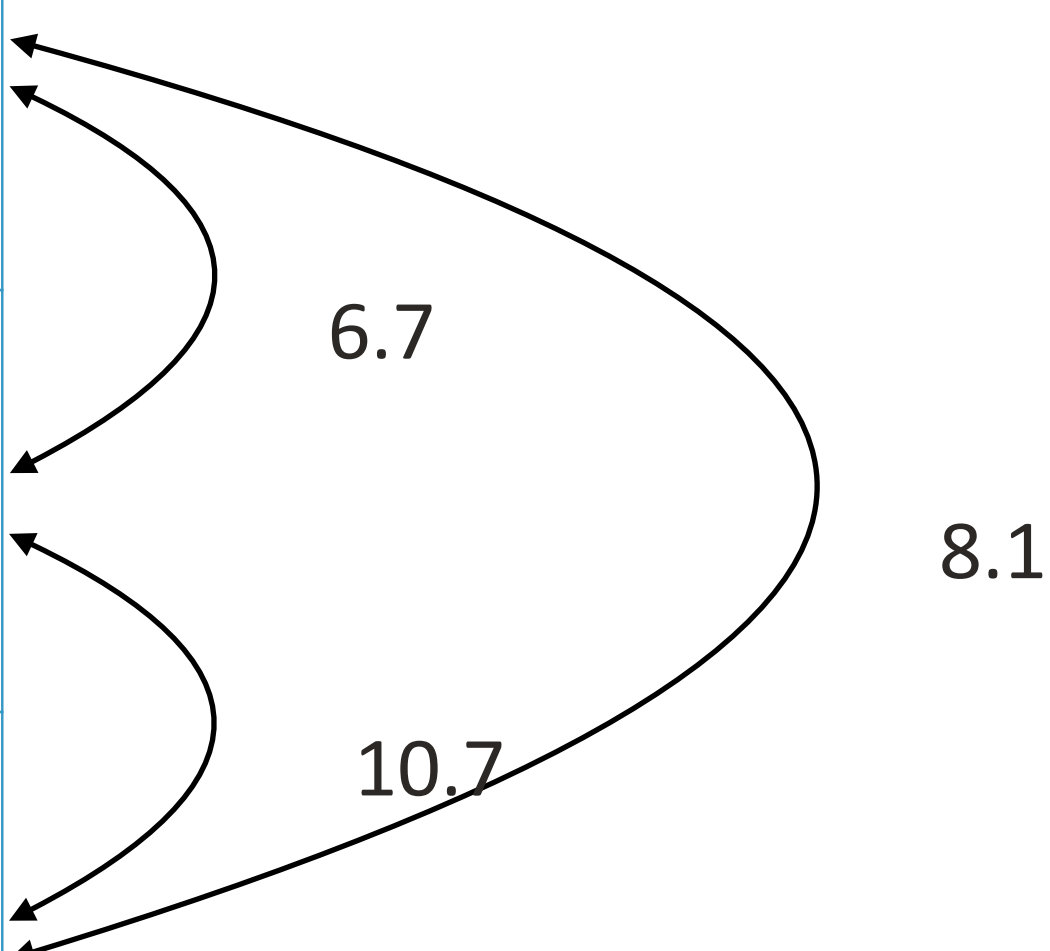
Dept1 to Dept3: ... = **5.4**




# Euclidean Distance calculations

```
def dist(x,y):  
    return np.sqrt(np.sum( (x-y)**2 ))  
  
> mat = np.array([[ 6,6,3 ], [1,2,1], [8,1,9]])  
> dist(mat[0], mat[1])  
6.7082039324993694  
  
> dist(mat[1], mat[2])  
10.677078252031311  
  
> dist(mat[0], mat[2])  
8.0622577482985491
```


# Which Departments are Similar?

	Malware events	Phishing	Open Tickets	
Dept1	6	6	3	
Dept2	1	2	1	
Dept3	8	1	9	


# Clustering...



# Clustering...



# Clustering...




# K-Means


Before starting, pick the number of clusters,  $K$

1. Pick  $K$  random centroids within data range
2. Assign each data point to the nearest centroid
3. Move centroid to center of assigned points
4. Repeat steps 2 and 3 until centroid stops shifting

# Step 1: Pick 3 random centroids within data range




# Step 2: Assign each data point to the nearest centroid (1)







# Step 3: Move centroid to center of assigned points (1)




# Step 2: Assign each data point to the nearest centroid (2)




# Step 3: Move centroid to center of assigned points (2)




# Step 2: Assign each data point to the nearest centroid (3)




# Step 3: Move centroid to center of assigned points (3)




# Step 2: Assign each data point to the nearest centroid (4)



# Step 3: Move centroid to center of assigned points (4)




# Step 2: Assign each data point to the nearest centroid (5)

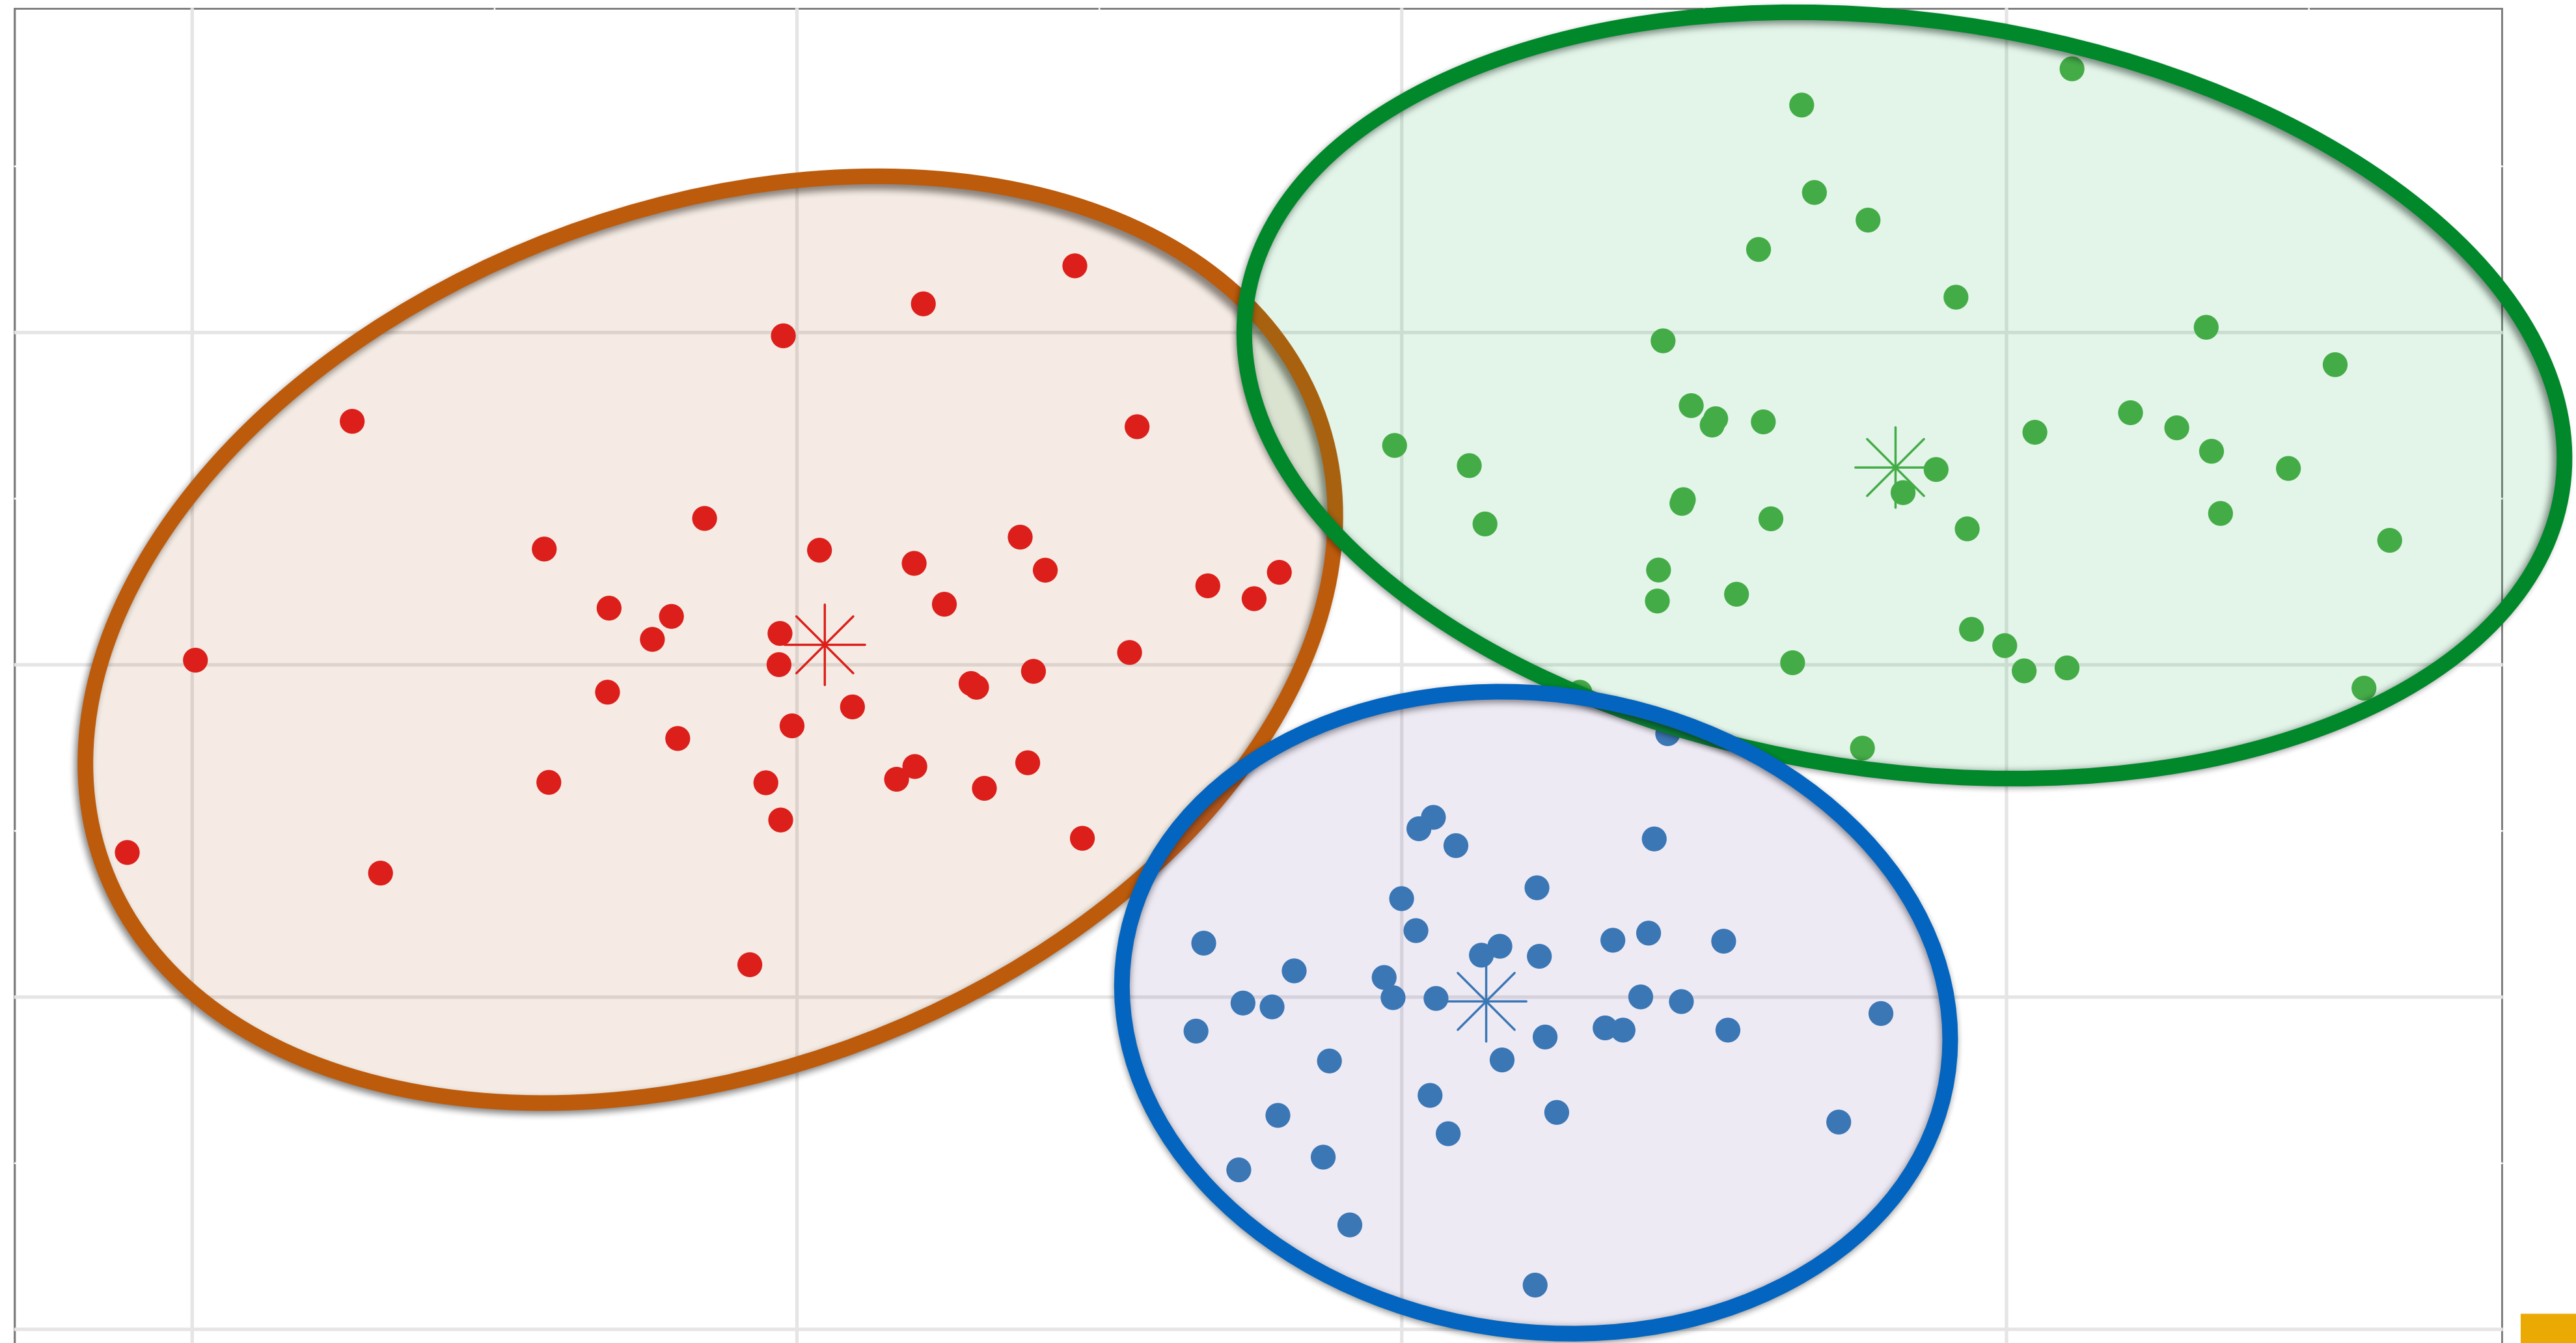




# Step 3: Move centroid to center of assigned points (5)



# Step 3: Move centroid to center of assigned points (5)



# K-Means: Got a problem with it?

Before starting, pick the number of clusters,  $K$


1. Pick  $K$  random centroids within data range
2. Assign each data point to the nearest centroid
3. Move centroid to center of assigned points
4. Repeat steps 2 and 3 until centroid stops shifting

# K-Means: Got a problem with it?


Before starting, pick the number of clusters, K *Subjective*

1. Pick K random centroids within data range *Not Repeatable*
2. Assign each data point to the nearest centroid *Sensitive to Scale*
3. Move centroid to center of assigned points
4. Repeat steps 2 and 3 until centroid stops shifting


# How many clusters?




# Random Start...



# Random Start...




# Random Start...





# K-Means




“...it’s too easy to throw k-means on your data, and nevertheless get a result out (that is pretty much random, but you won't notice).”


— Anony-Mousse

<http://stats.stackexchange.com/questions/133656/how-to-understand-the-drawbacks-of-k-means>

# Which one is correct?



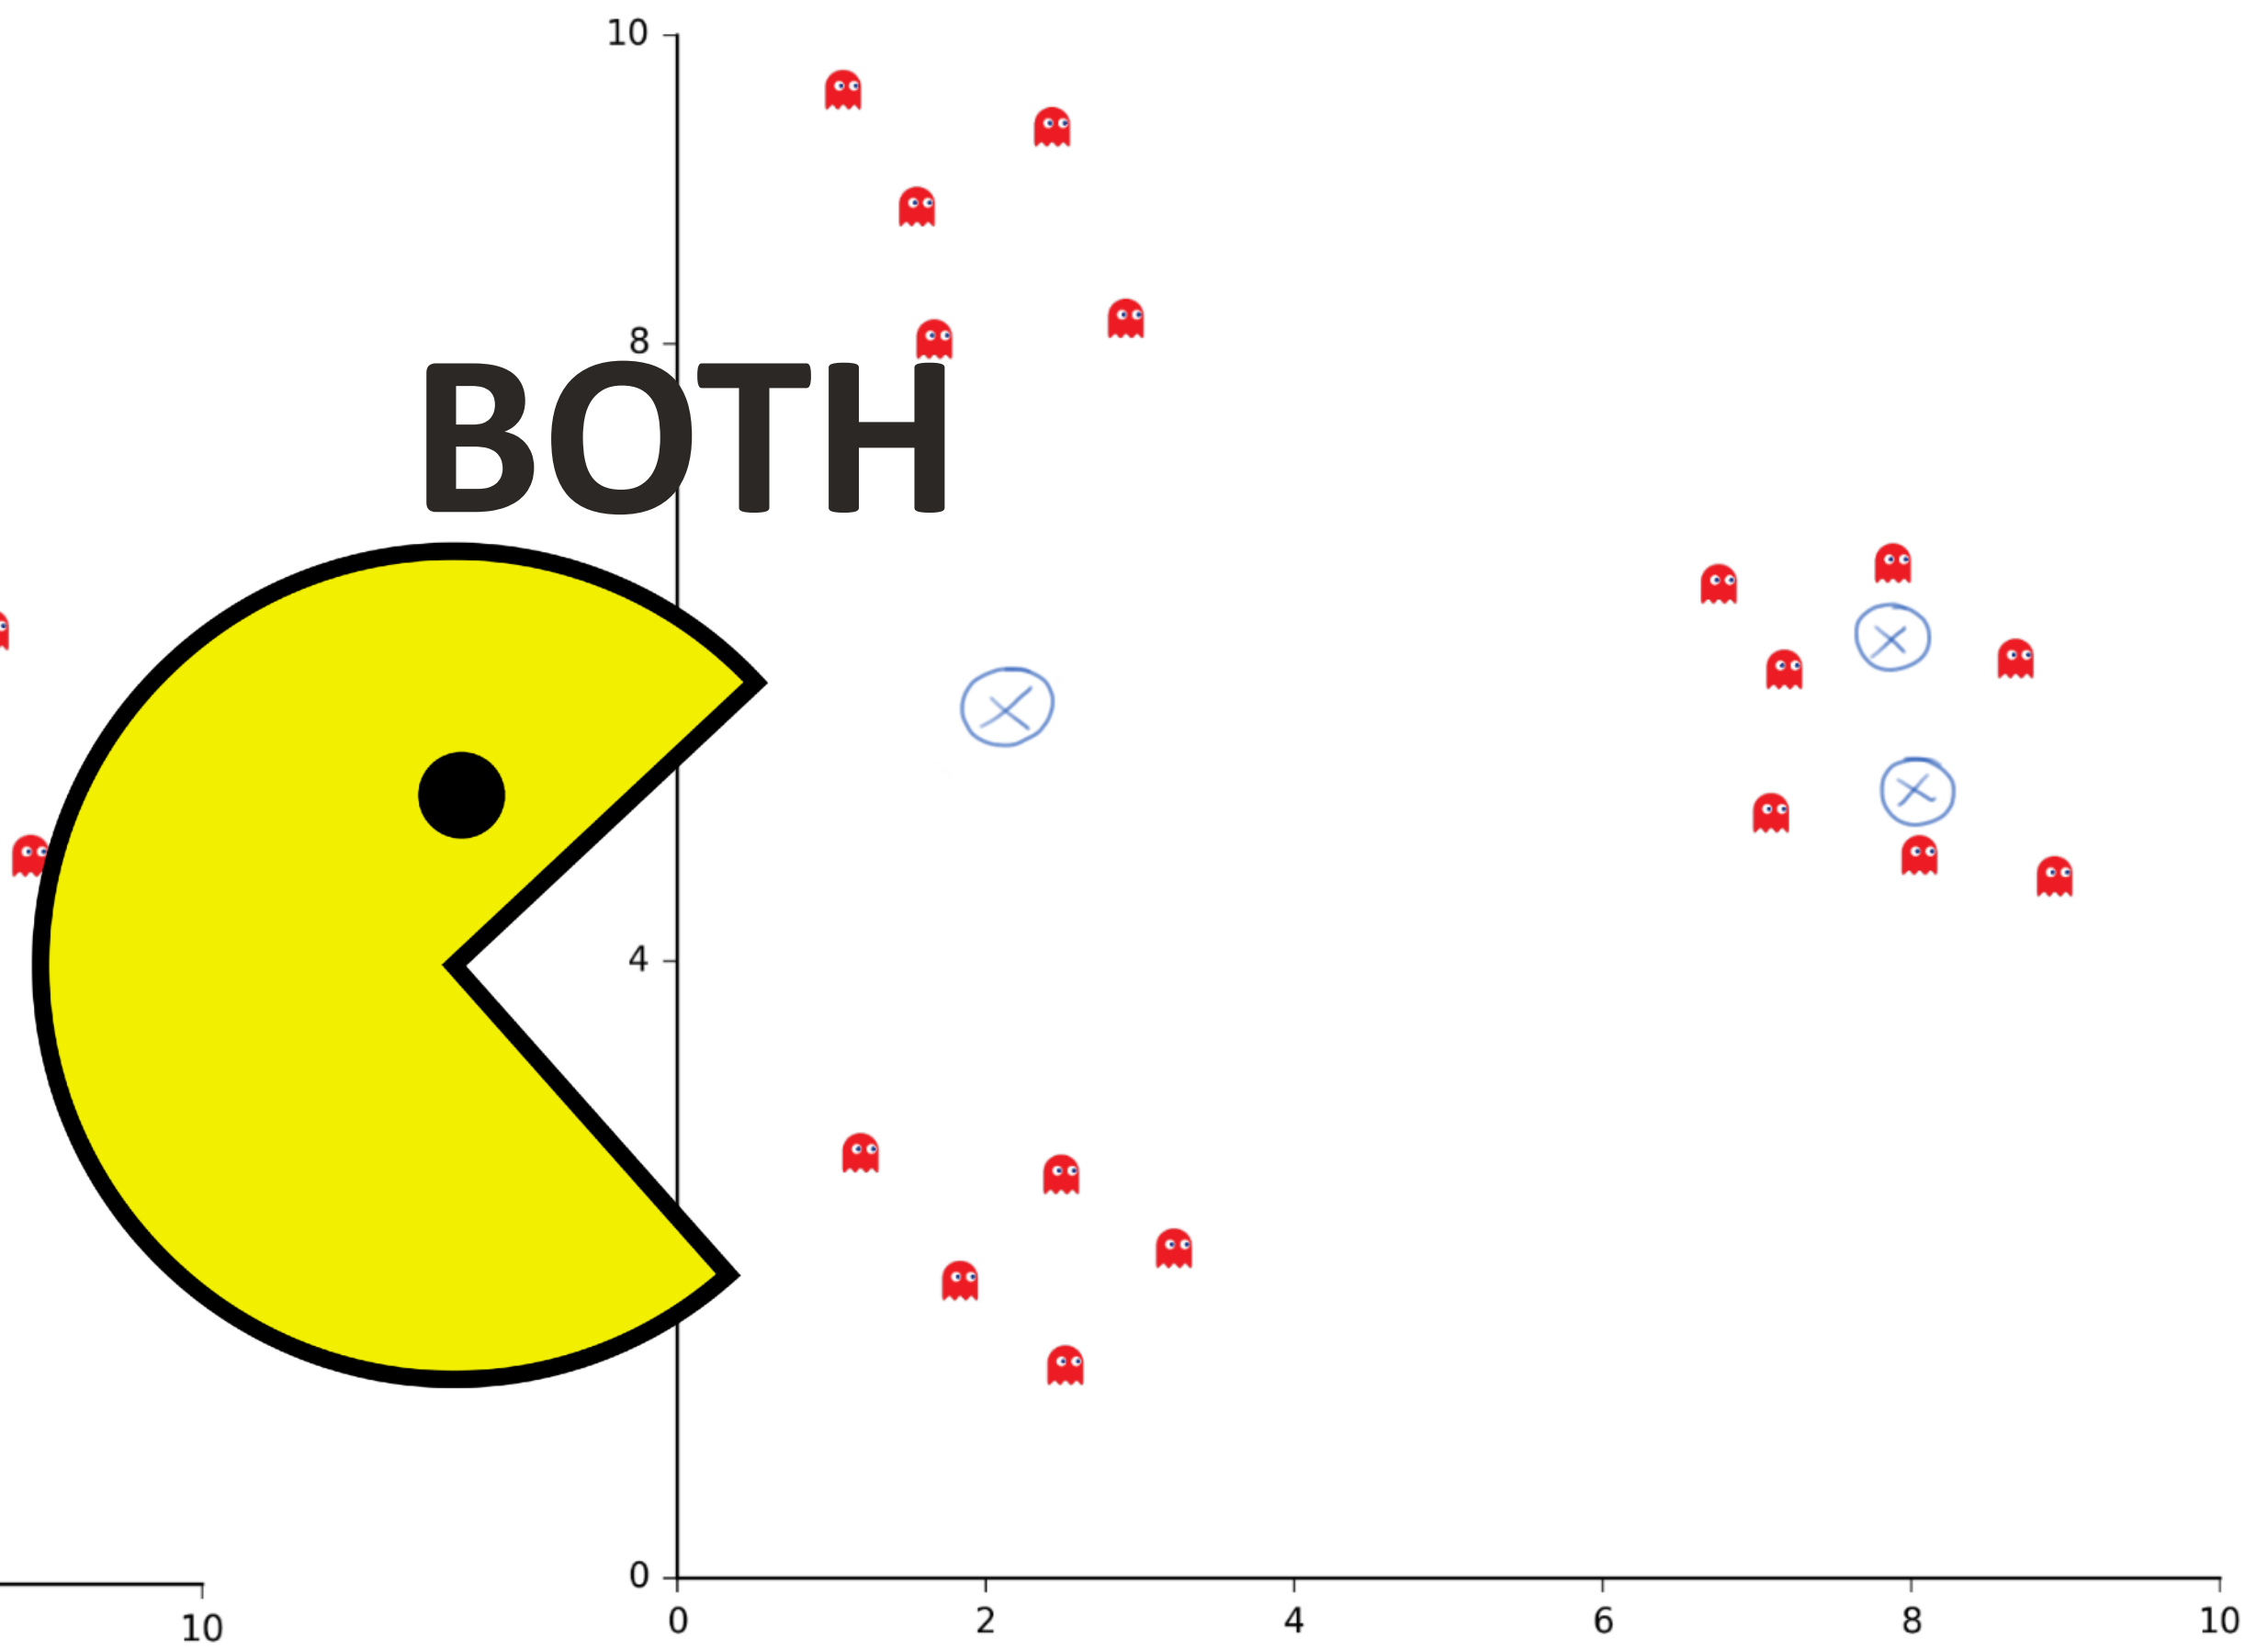

A




B



# Which one is correct?



# Pain of optimization...Being stuck at sub-optimal local minimum...




Initial guess matters!  
Same outcome cannot be guaranteed

# K-Means in practice (Python version)

```
#Import from Scikit-learn  
from sklearn.cluster import KMeans
```


```
kmeans = KMeans(n_clusters=2)  
kmeans.fit(data)
```

```
centroids = kmeans.cluster_centers_  
labels = kmeans.labels_f
```



# The Dataset

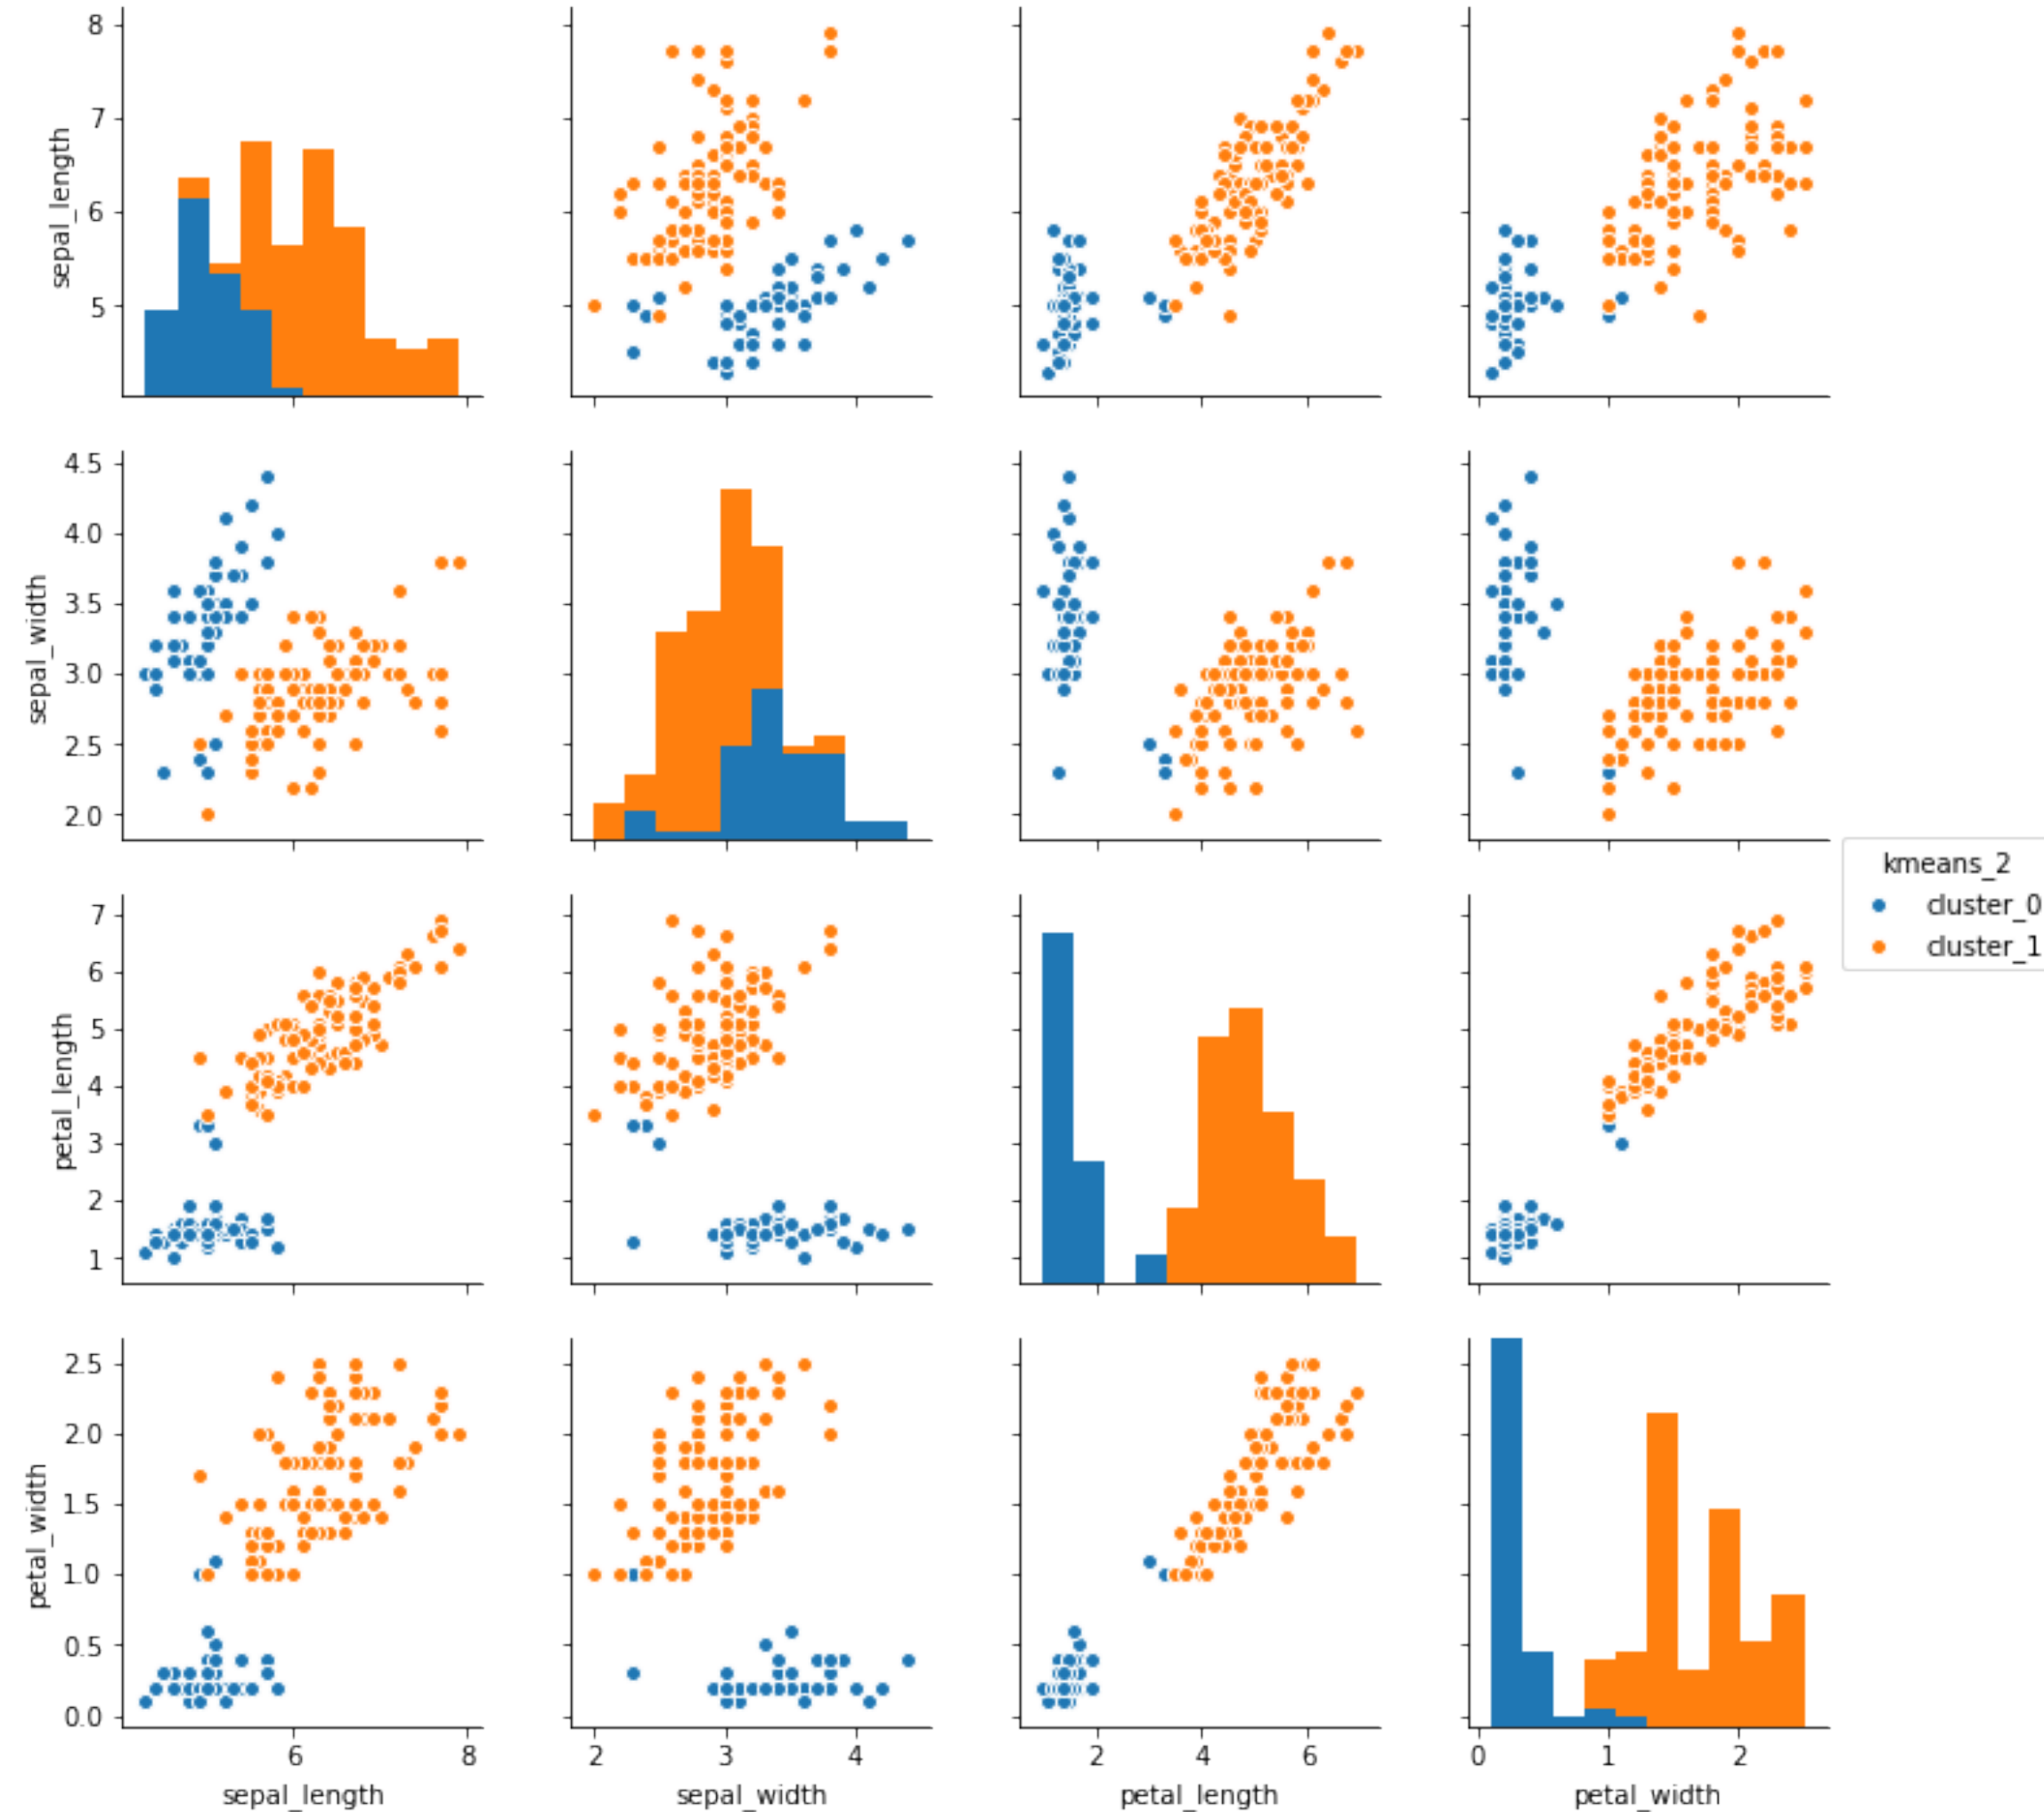
```
sns.pairplot(<data>)
```



# K-Means Clustering

```
kmeans = KMeans(n_clusters=2, n_init=2)  
kmeans.fit( <data> )
```

# K-Means Clustering




```
sns.pairplot(<data>, hue="kmeans_2")
```



# K-Means Clustering

```
sns.pairplot(<data>, x_vars="col_1", y_vars="col_2", hue="kmeans_2", size=6)  
plt.scatter(<cluster_centers>, <col_2>, linewidths=3, marker='x', s=200,  
c='black' )
```



**K-Means is affected by the scale  
of every feature.**

# Feature Scaling

For k-means clustering, features must be scaled to the same ranges of values to contribute "equally" to the euclidean distance calculation.

Each row is transformed per-column by:

- Subtracting from the element in each row the mean for each feature (column) and then taking this value and
- Dividing by that feature's (column's) standard deviation.

# Feature Scaling

```
# center and scale the data  
scaler = StandardScaler()
```

```
raw_data_scaled = scaler.fit_transform( <data> )
```


```
data_scaled = pd.DataFrame( raw_data_scaled,  
columns=features )
```

# Feature Scaling

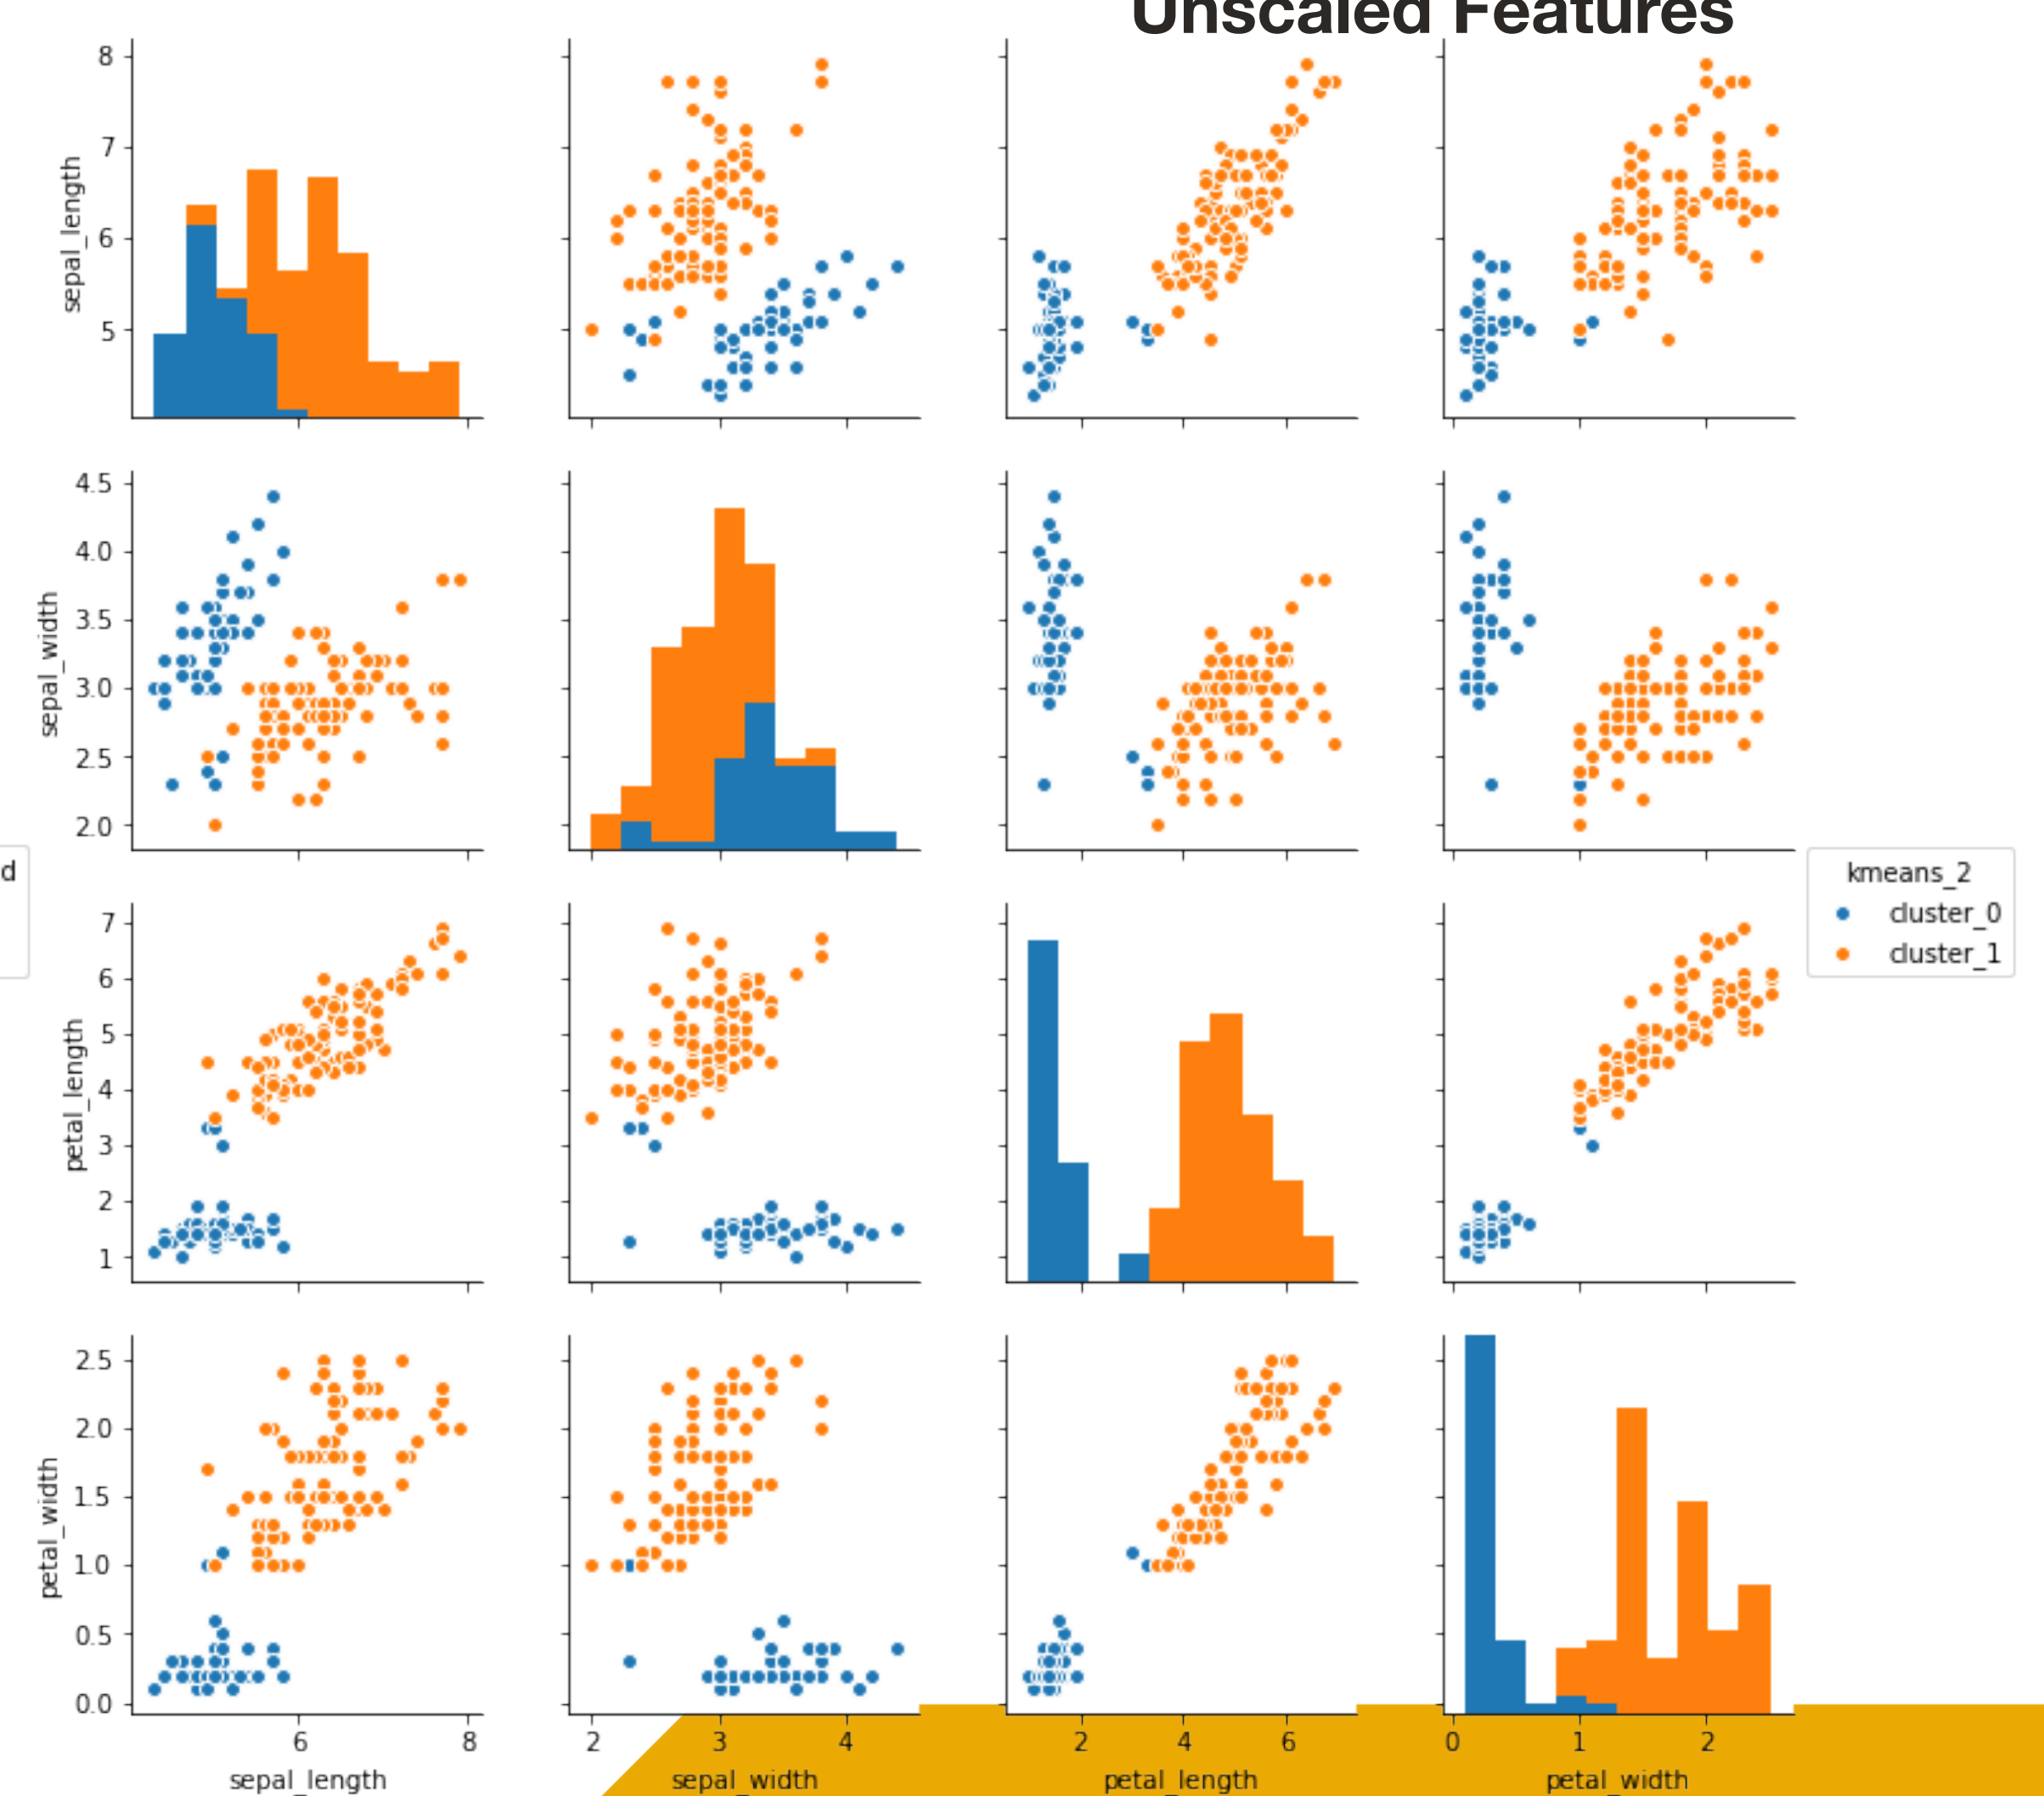
```
# K-means on scaled data  
km = KMeans(n_clusters=2, n_init=2)  
km.fit( <scaled_data> )
```

# Feature Scaling

Scaled Features



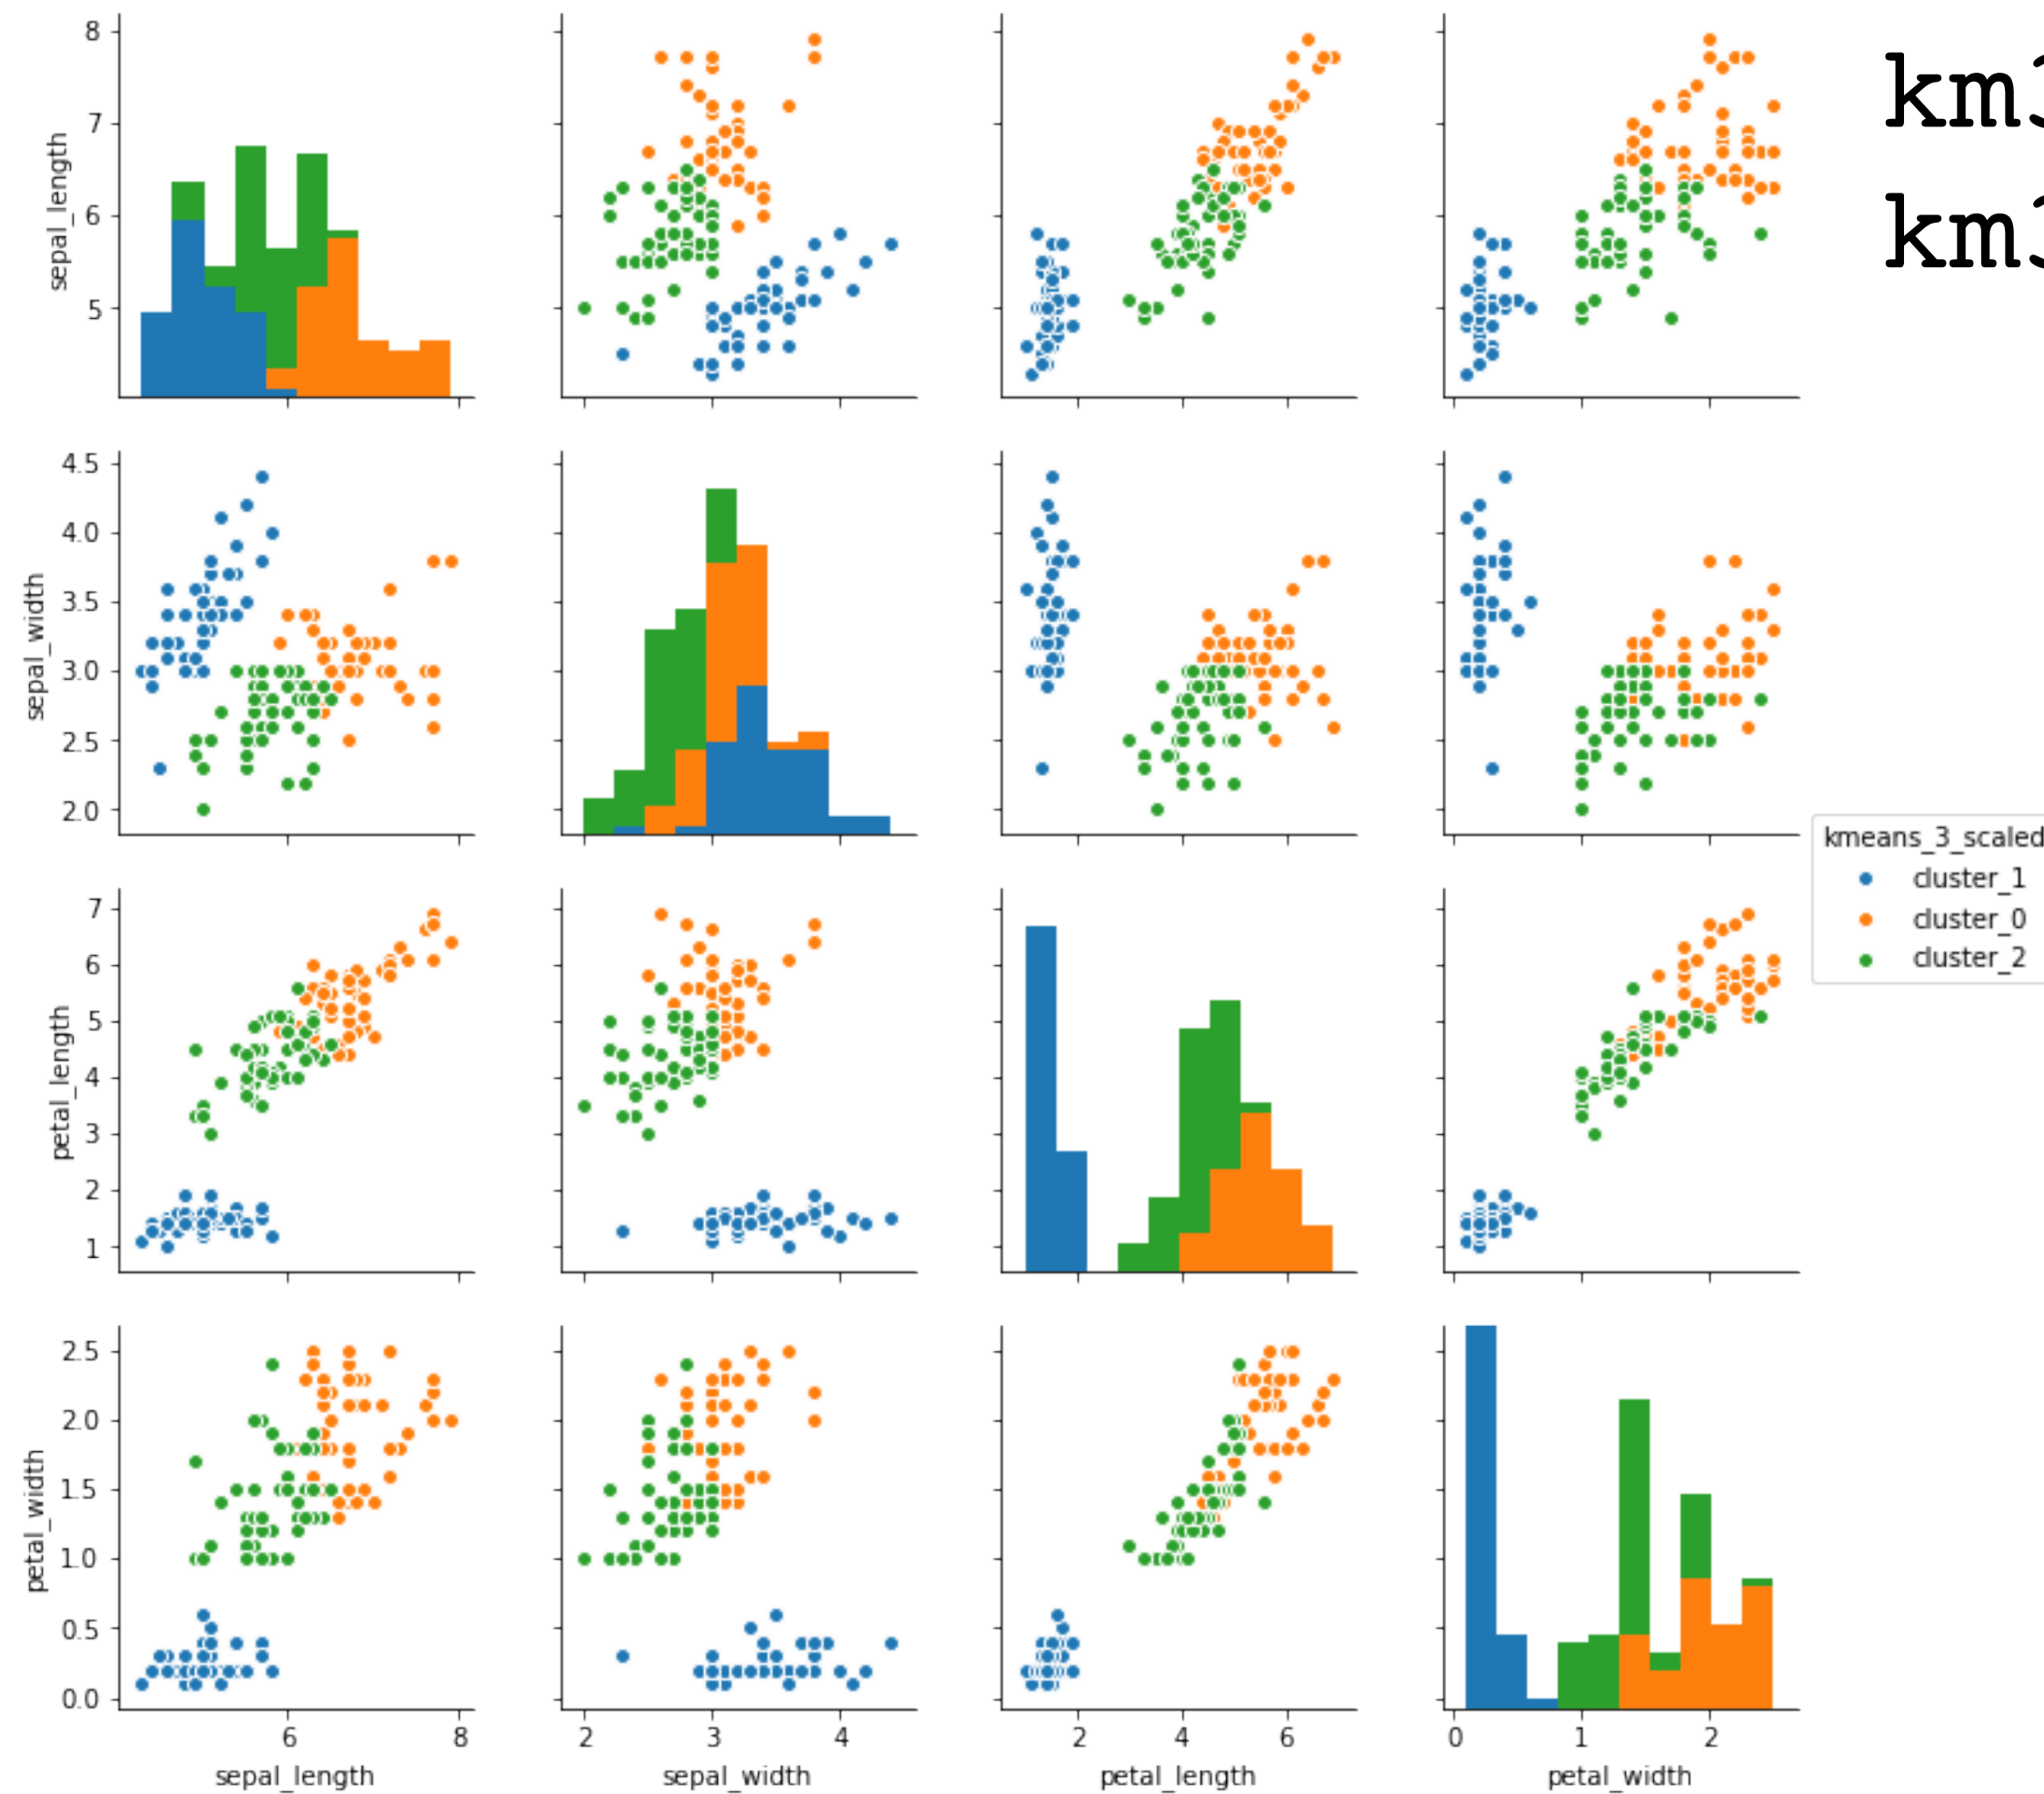
Unscaled Features





# More Clusters


```
km3 = KMeans(n_clusters=3, n_init=2)  
km3.fit(scaled_data)
```



# Outlier Detection

```
clf = svm.OneClassSVM( tol=0.001, nu=0.1)
clf.fit(X)
target_pred_outliers=clf.predict(X)
```

Delete n% of “outlier data”,  
here ~10%





# Evaluating your Model

The Silhouette Coefficient is a common metric for evaluating clustering "performance" in situations when the "true" cluster assignments are not known.

b = mean distance to next nearest cluster


a = mean distance to other points in cluster

$$\text{silhouette\_coeff} = (b - a) / \max(a, b)$$

# Evaluating your Model

```
k_range = range(2,16)
scores = []
for k in k_range:
    km_ss = KMeans(n_clusters=k, random_state=1)
    km_ss.fit(iris_data_scaled)
    scores.append(silhouette_score(<data>,
km_ss.labels_))
```


# Evaluating your Model



# Evaluating your Model

```
from yellowbrick.cluster import KElbowVisualizer  
visualizer = KElbowVisualizer(KMeans(), k=(4,12))
```


```
visualizer.fit(X)  
visualizer.poof()
```



# Evaluating your Model

```
from yellowbrick.cluster import SilhouetteVisualizer  
model = MiniBatchKMeans(6)  
visualizer = SilhouetteVisualizer(model)
```

```
visualizer.fit(X)  
visualizer.poof()
```



# DBSCAN

DBSCAN stands for **D**ensity-**B**ased **S**patial **C**lustering of **A**pplications with **N**oise.

Whereas K-means does not care about the density of data, DBSCAN does, under the assumption that regions of high density in your data should be treated as clusters.

# DBSCAN

DBSCAN does not allow you to specify how many clusters you want. Instead, you specify 2 parameters:

- **e (epsilon)**: This is the maximum distance between two points to allow them to be neighbors
- **min\_samples**: The number of neighbors a given point is allowed to have to be able to be part of a cluster

Any points that don't satisfy the criteria of being close enough to other points are labeled outliers and all fall into a single "cluster" (their cluster label by default is -1).

# DBSCAN

DBSCAN works as follows:

1. Choose an arbitrary starting point in your dataset that has not been seen.
2. Retrieve this point's  $\epsilon$ -neighborhood (all points that are within a distance  $\epsilon$  from it), and if it contains at least **\*min\_samples**, a cluster is started.
3. Otherwise, the point is labeled as an outlier (-1). Note: This point might later be found in a sufficiently sized  $\epsilon$ -environment of a different point and hence be made part of a cluster.
4. If a point is found to be a dense part of a cluster, its  $\epsilon$ -neighborhood is also part of that cluster. All points that are found within the  $\epsilon$ -neighborhood are added, as is their own  $\epsilon$ -neighborhood when they are also dense.
5. Continue until the density-connected cluster is completely found.
6. Find a new unvisited point to process and repeat.




# DBSCAN

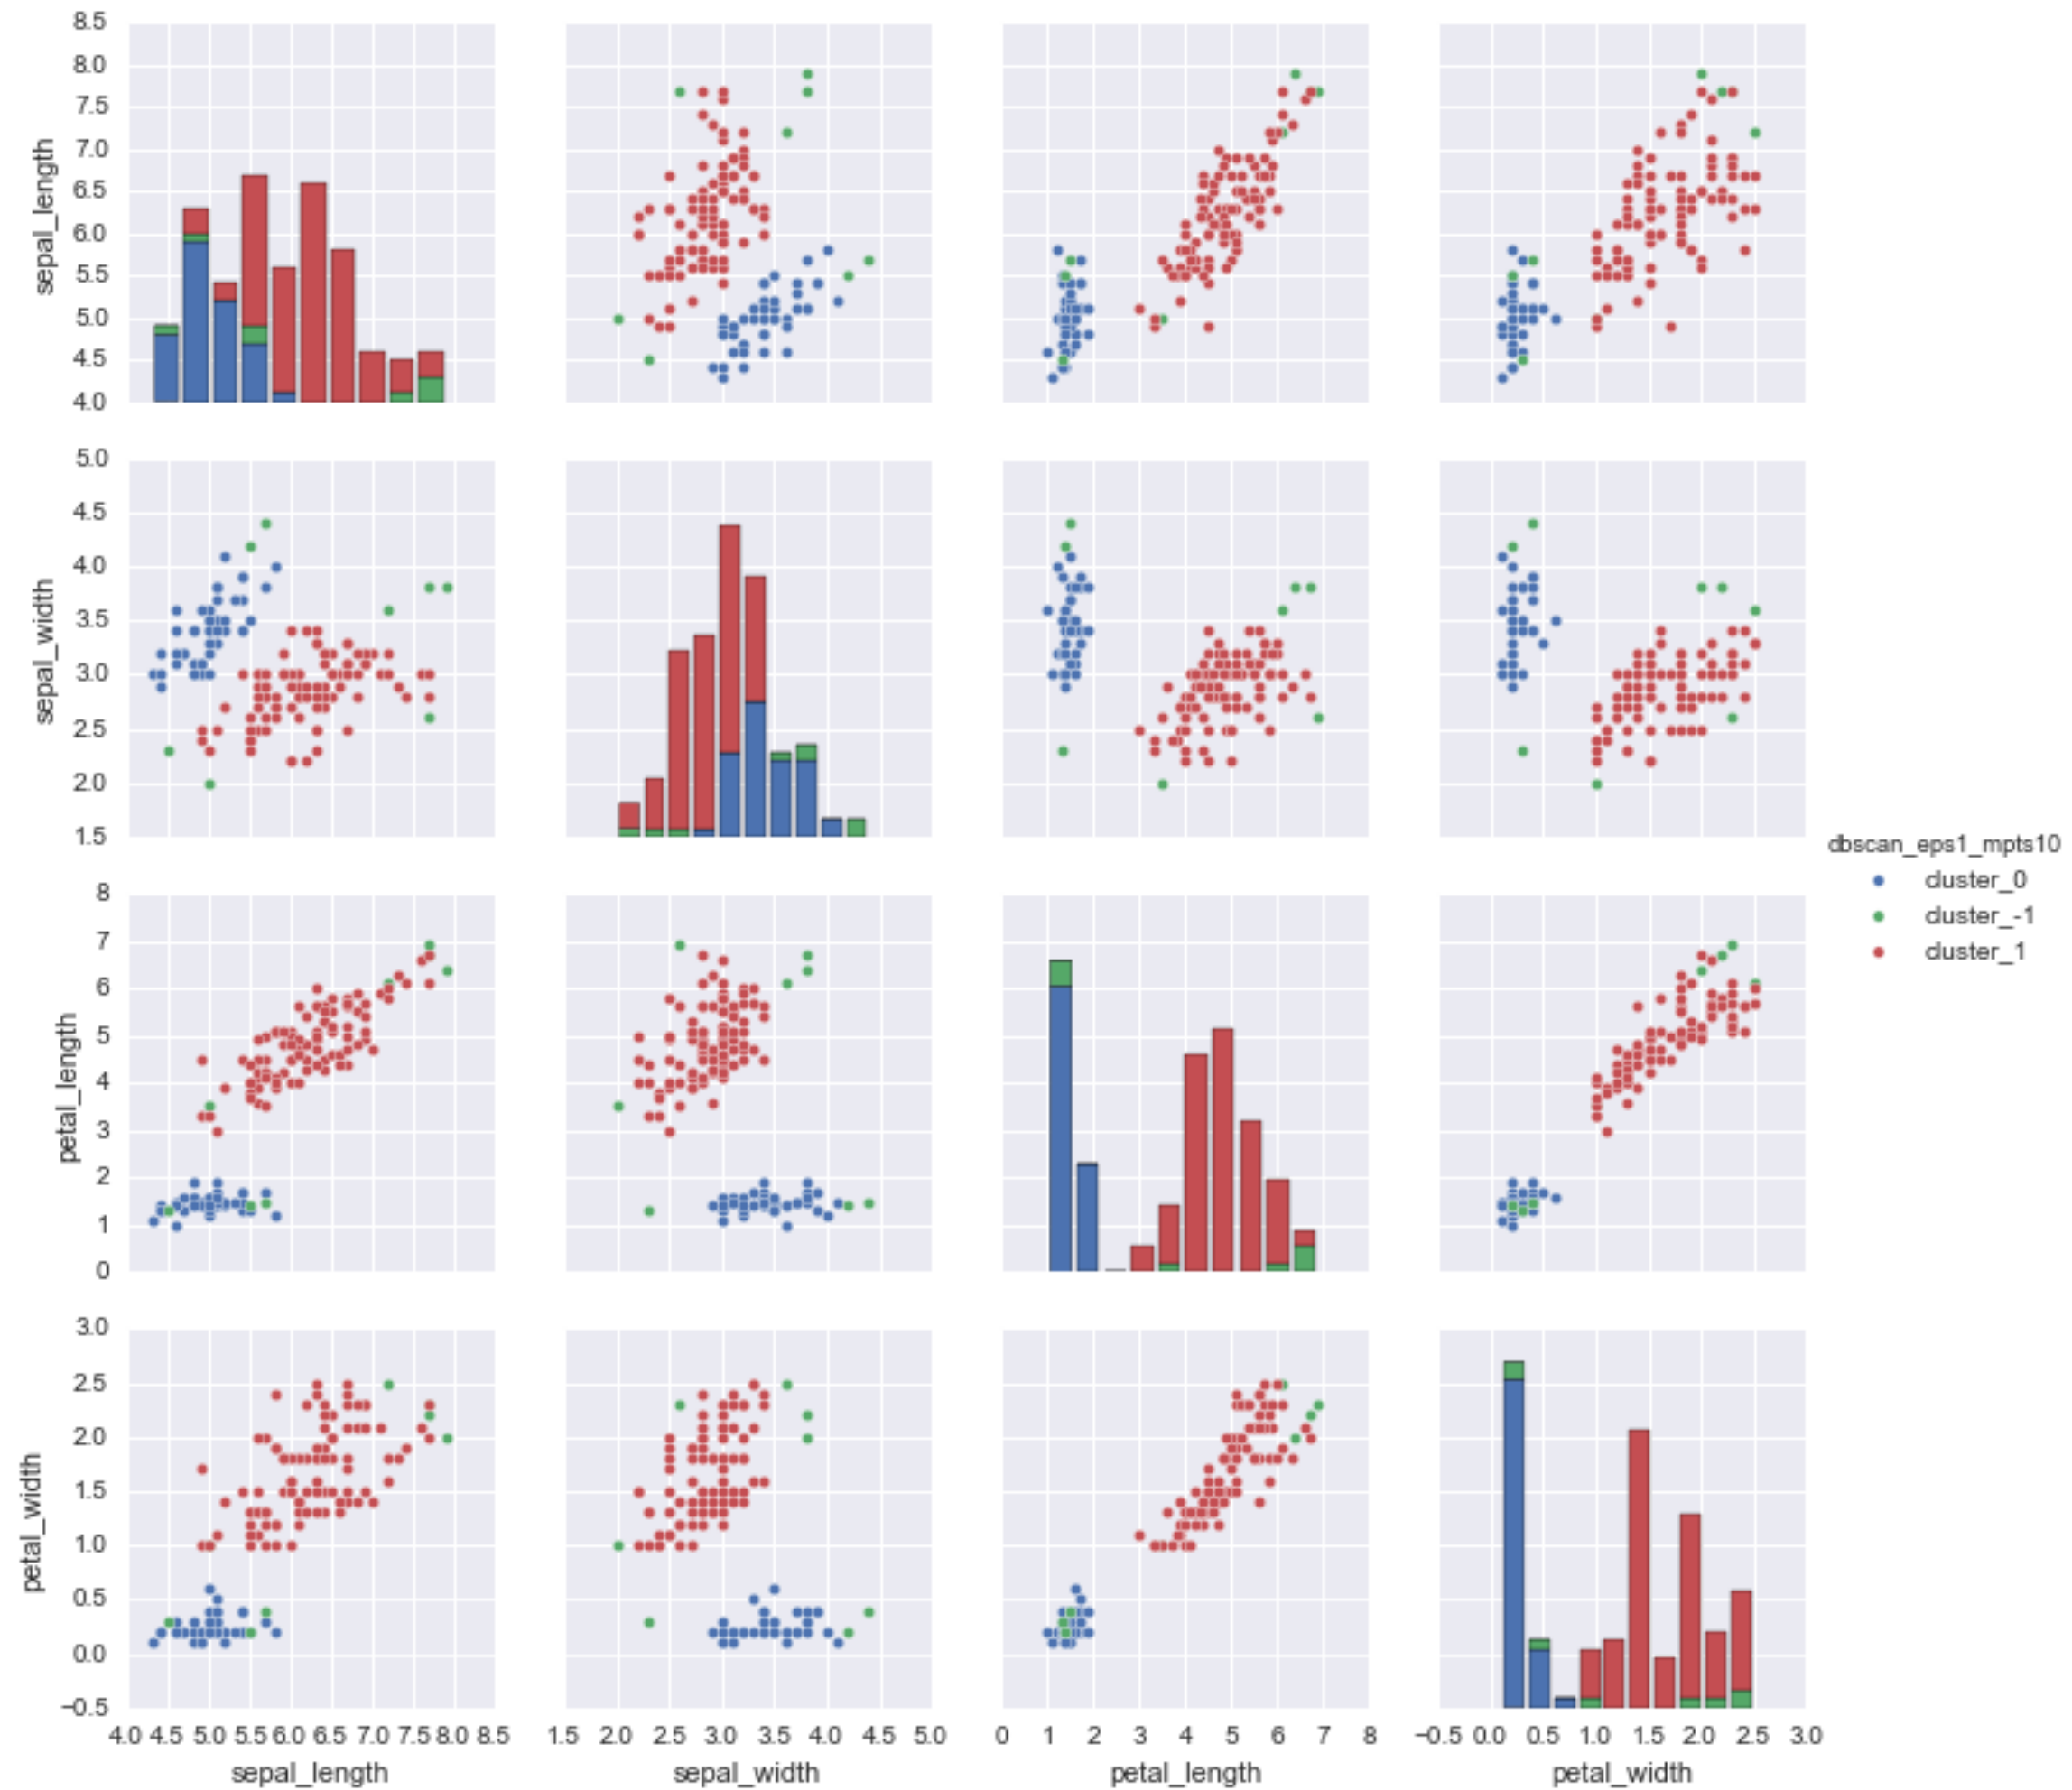
```
db = DBSCAN(eps=1, min_samples=3)  
db.fit(<scaled_data>)
```

# DBSCAN

```
data_no_names['dbscan_eps1_mpts3'] = [ "cluster_" + str(label) for label in db.labels_ ]  
sns.pairplot(data_no_names,hue="dbscan_eps1_mpts3")
```




# DBSCAN



```
db2 = DBSCAN(eps=1, min_samples=10)
db2.fit(data_scaled)
```


# DBSCAN



```
db2 = DBSCAN(eps=2, min_samples=3)
db2.fit(iris_data_scaled)
```

# DBSCAN

```
db2 = DBSCAN(eps=0.6, min_samples=3)  
db2.fit(iris_data_scaled)
```



# **Other uses of unsupervised learning: Dimensionality Reduction**



# Other uses of unsupervised learning: Dimensionality Reduction

- Dimensionality reduction reduces the number of features in a dataset without losing the information in a data set
- Common technique: **Principal Component Analysis (PCA)**

# Other uses of unsupervised learning: Dimensionality Reduction

- PCA attempts to combine the highly correlated features and represent this data with a smaller number of linearly uncorrelated features.
- The algorithm keeps performing this correlation reduction, finding the directions of maximum variance in the original high dimensional data and projecting them onto a smaller dimensional space.
- These newly derived components are known as principal components.



# PCA Example

```
from sklearn.decomposition import PCA
```

```
n_components = 784
```

```
whiten = False
```

```
pca = PCA(n_components=n_components, whiten=whiten)
```

```
features_train_PCA = pca.fit_transform(features_train_PCA)
```

```
features_train_PCA = pd.DataFrame(data= features_train_PCA,  
index=train_index)
```

```
print("Variance Explained by all 784 principal components: ", \  
      sum(pca.explained_variance_ratio_))
```

The variance of all features is 1.0 (everything)

# PCA Example

```
Variance Captured by First 10 Principal Components: [0.48876238]
Variance Captured by First 20 Principal Components: [0.64398025]
Variance Captured by First 50 Principal Components: [0.8248609]
Variance Captured by First 100 Principal Components: [0.91465857]
Variance Captured by First 200 Principal Components: [0.96650076]
Variance Captured by First 300 Principal Components: [0.9862489]
```

**Bottom line:** 98% of the information is contained in 300 features. You can therefore, reduce the size of your dataset by half, theoretically without significantly reducing performance.

[https://github.com/aapatel09/hands-on-unsupervised-learning/blob/master/03\\_dimensionality\\_reduction.ipynb](https://github.com/aapatel09/hands-on-unsupervised-learning/blob/master/03_dimensionality_reduction.ipynb)

# Questions?

# In Class Exercise

Please complete Worksheet 6.0: Clustering